

Lecture 2 Performance

Fall 2016

Pejman Lotfi-Kamran



Adapted from slides originally developed by Profs. Falsafi, Hill, Hoe, Lipasti, Shen, Smith, Sohi, and Vijaykumar of Carnegie Mellon University, EPFL, Purdue University, and University of Wisconsin.

Lecture 2
Slide 1

Where Are We?

| Fr | Sa | Su | Mo | Tu |
|----|--------------|----|--------------|----|
| | 27-Shahrivar | | 29-Shahrivar | |
| | 3-Mehr | | 5-Mehr | |
| | 10-Mehr | | 12-Mehr | |
| | 17-Mehr | | 19-Mehr | |
| | 24-Mehr | | 26-Mehr | |
| | 1-Aban | | 3-Aban | |
| | 8-Aban | | 10-Aban | |
| | 15-Aban | | 17-Aban | |
| | 22-Aban | | 24-Aban | |
| | 29-Aban | | 1-Azar | |
| | 6-Azar | | 8-Azar | |
| | 13-Azar | | 15-Azar | |
| | 20-Azar | | 22-Azar | |
| | 27-Azar | | 29-Azar | |
| | 4-Dey | | 6-Dey | |

This Lecture

- Performance measurement

Next Lecture:

- Basic caches

Lecture 2
Slide 2

Metrics of Performance

Time (latency)

- elapsed time vs. processor time

Rate (bandwidth or throughput)

- performance = rate = work per time

Distinction is sometimes blurred

- consider batched vs. interactive processing
- consider overlapped vs. non-overlapped processing
- may require conflicting optimizations

What is the most popular metric of performance?

Lecture 2
Slide 3

The “Iron Law” of Processor Performance

$$\text{Processor Performance} = \frac{\text{Time}}{\text{Program}}$$

$$= \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Cycles}}{\text{Instruction}} \times \frac{\text{Time}}{\text{Cycle}}$$

(code size) (CPI) (cycle time)

Architecture --> Implementation --> Realization

Compiler Designer

Processor Designer

Chip Designer

Lecture 2
Slide 4

MIPS

MIPS - Millions of Instructions per Second

$$\text{MIPS} = \frac{\text{\# of instructions}}{\text{benchmark}} \times \frac{\text{benchmark}}{\text{total run time}}$$

1,000,000

When comparing two machines (A, B) with the same **instruction set**, MIPS is a fair comparison (*sometimes...*)

But, MIPS can be a “*meaningless indicator of performance...*”

- instruction sets are not equivalent
- different programs use a different instruction mix
- instruction count is not a reliable indicator of work
 - some optimizations add instructions
 - instructions have varying work

Lecture 2
Slide 5

MIPS (Cont.)

Example:

- Machine A has a special instruction for performing square root
 - It takes 100 cycles to execute
- Machine B doesn't have the special instruction
 - must perform square root in software using simple instructions
 - e.g, Add, Mult, Shift each take 1 cycle to execute
- Machine A: 1/100 MIPS = 0.01 MIPS
- Machine B: 1 MIPS

How about Relative MIPS = $(\text{time}_{\text{reference}} / \text{time}_{\text{new}}) \times \text{MIPS}_{\text{reference}}$?

Lecture 2
Slide 6

MFLOPS

$$\text{MFLOPS} = (\text{FP ops/program}) \times (\text{program/time}) \times 10^{-6}$$

Popular in scientific computing

- There was a time when FP ops were much much slower than regular instructions (i.e., off-chip, sequential execution)

Not great for “predicting” performance because it

- ignores other instructions (e.g., load/store)
- not all FP ops are created equally
- depends on how FP-intensive program is



Beware of “peak” MFLOPS!

Lecture 2
Slide 7

Normalized MFLOPS

Normalized FP: give canonical # FP ops to prog

$$\text{Normalized MFLOPS} = (\text{\# canonical FP ops/time}) \times 10^{-6}$$

Not all machines have the same FP ops

- Cray does not implement divide
- Motorola has SQRT, SIN, and COS

Not all FP ops do the same amount of work

- adds usually faster than divide



Used to compare performance, essentially a measure of execution time.

Lecture 2
Slide 8

Comparing Performance

Often, we want to compare the performance of different machines or different programs. Why?

- help architects understand which is “better”
- give marketing a “silver bullet” for the press release
- help customers understand why they should buy <my machine>

Lecture 2
Slide 9

Comparing Performance (Cont.)

If machine A is 50% slower than B, and $\text{time}_A = 1.0\text{s}$, does that mean

case 1: $\text{time}_B = 0.5\text{s}$ since $\text{time}_B/\text{time}_A = 0.5$

case 2: $\text{time}_B = 0.666\text{s}$ since $\text{time}_A/\text{time}_B = 1.5$

What if I say machine B is 50% faster than A?

If machine A is 1000% faster than B, and $\text{time}_A = 1.0\text{s}$, does that mean

case 1: $\text{time}_B = 10.0\text{s}$ A is 10 times faster

case 2: $\text{time}_B = 11.0\text{s}$ A is 11 times faster

Lecture 2
Slide 10

By Definition

Machine A is n times faster than machine B iff
 $\text{perf}(A)/\text{perf}(B) = \text{time}(B)/\text{time}(A) = n$

Machine A is $x\%$ faster than machine B iff
 $\text{perf}(A)/\text{perf}(B) = \text{time}(B)/\text{time}(A) = 1 + x/100$

E.g., A 10s, B 15s

$15/10 = 1.5 \Rightarrow$ A is 1.5 times faster than B

$15/10 = 1 + 50/100 \Rightarrow$ A is 50% faster than B

Remember to compare “performance”

Lecture 2
Slide 11

Performance vs. Execution Time

Often, we use the phrase “*X is faster than Y*”

- Means the response time or execution time is lower on X than it is on Y
- Mathematically, “X is N times faster than Y” means

$$\frac{\text{Execution Time}_Y}{\text{Execution Time}_X} = N$$

$$\frac{\text{Execution Time}_Y}{\text{Execution Time}_X} = N = \frac{1/\text{Performance}_Y}{1/\text{Performance}_X} = \frac{\text{Performance}_X}{\text{Performance}_Y}$$

Performance and Execution time are *reciprocals*

- Increasing performance decreases execution time

Lecture 2
Slide 12

Reasons to Compare Performance

To make a decision

- customers: which machine to buy
- designers: which optimization to include vs. leave out

Important to have “representative” performance

- ~~Time/prog~~ = insts/prog x cycles/inst x sec/cycle

Lecture 2
Slide 13

Type of Benchmarks

Mixes

- instruction frequency of occurrence; calculate

Kernels

- “representative” program fragments
- good for focusing on individual features not big picture

Synthetic benchmarks

- programs intended to give specific mix
- ignore dependences
- maybe ok for non-pipelined, non-cached, w/o optimizing compilers
- questionable validity
- simple to measure and report

Lecture 2
Slide 14

Type of Benchmarks (cont.)

Real programs

- representative of real workload
- more accurate way to characterize performance
- requires considerable work

But, nothing beats testing out your bread&butter application!!

Lecture 2
Slide 15

Mix Example

Gibson Mix, developed in 1950's at IBM

| | | | |
|---------------|------|------------|-----|
| load/store | 31% | branch | 17% |
| fixed add/sub | 6% | compare | 4% |
| float add/sub | 7% | float mult | 4% |
| float div | 2% | fixed mult | 1% |
| fixed div | < 1% | shifts | 4% |
| logical | 2% | | |

Generally speaking, these numbers are still valid today.

But, machine behaviors are too complicated

Lecture 2
Slide 16

Kernel Example

```

Inner product
DO L = 1, LP
  Q = 0.0
DO K = 1, N
  Q = Q + Z(K)*X(K)

```

Lecture 2
Slide 17

Synthetic Benchmark Example

```

Dhrystone, Whetstone
X = 1.0
Y = 1.0
Z = 1.0
DO I = 1, N8
  CALL P3(X,Y,Z)
SUBROUTINE P3(X,Y,Z)
  X1 = X
  Y1 = Y
  X1 = T * (X1 - Y1)
  Y1 = T * (X1 + Y1)
  Z = (X1 + Y1)/T2
  RETURN

```

Intended to measure procedure call and return

How about $X = \text{SQRT}(\text{EXP}(\text{ALOG}(X)/T1))$?

Lecture 2
Slide 18

SPEC

SPEC - The System Performance Evaluation Cooperative (SPEC)

- ❑ founded in 1988 by a small number of workstation vendors who realized that the marketplace was in desperate need of realistic, standardize performance tests.
- ❑ Grown to become successful performance standardization bodies with more than 40 member companies. <http://www.spec.org>

SPEC's Philosophy

- ❑ The goal of SPEC is to ensure that the marketplace has a fair and useful set of metrics to differentiate candidate systems.
- ❑ The basic SPEC methodology is to provide the benchmark with a standardized suite of source code based upon existing applications

Benchmarks come in many different flavors

- ❑ CPUINT && CPUFP 2006
- ❑ JVM08
- ❑ SFS – NFS benchmarks
- ❑ WEB09

Lecture 2
Slide 19

SPEC CPUInt 2006 Benchmarks

| | | |
|----------------|-----|---------------------------------|
| 400.perlbench | C | Programming language |
| 401.bzip2 | C | Compression |
| 403.gcc | C | C Programming Language Compiler |
| 429.mcf | C | Combinatorial Optimization |
| 445.gobmk | C | Artificial Intelligence: Go |
| 456.hmmcr | C | Search Gene Sequence |
| 458.sjeng | C | Artificial Intelligence: Chess |
| 462.libquantum | C | Physics / Quantum Computing |
| 464.h264ref | C | Video Compression |
| 471.omnetapp | C++ | Discrete Event Simulation |
| 473.astar | C++ | Path-Finding Algorithms |
| 483.xalanbmk | C++ | XML Processing |

Lecture 2
Slide 20

SPEC CPUFP 2006 Benchmarks

| | | |
|---------------|------------|------------------------------------|
| 410.bwaves | Fortran | Fluid Dynamics |
| 416.gamess | Fortran | Quantum Chemistry |
| 433.milc | C | Physics / Quantum Chromodynamics |
| 434.zeusmp | Fortran | Physics / CFD |
| 435.gromacs | C, Fortran | Bio Chemistry / Molecular Dynamics |
| 436.cactusADM | C, Fortran | Physics / General Relativity |
| 437.leslie3d | Fortran | Fluid Dynamics |
| 444.Namd | C++ | Biology / Molecular Dynamics |
| 447.dealII | C++ | Finite Element Analysis |
| 450.soplex | C++ | Linear Programming, Optimization |
| 453.povray | C++ | Image Ray-Tracing |
| 454.calculix | C, Fortran | Structural Mechanics |
| 459.GemsFDTD | Fortran | Computational Electromagnetics |
| 465.tonto | Fortran | Quantum Chemistry |
| 470.lbm | C | Fluid Dynamics |
| 481.wrt | C, Fortran | Weather |
| 482.sphinx3 | C | Speech Recognition |

Lecture 2
Slide 21

SPEC Performance Numbers

Geometric Mean of 12 (SpecINT) and 17 (SpecFP) Benchmarks

□ Performance measured against Sun UltraSparc II system at 296MHz










2006 Performance Numbers (www.spec.org)

| | AMD Opteron 2.4GHz | IBM Power 5 1.9GHz | Intel Itanium 2 1.5GHz | Intel Pentium 4 3.4GHz | SGI R14K 500MHz | Sun Ultra III Cu 1.2GHz |
|------|--------------------------|--------------------------|------------------------------|------------------------------|-----------------------|-------------------------------|
| Int* | 1346 | 1398 | 1408 | 1669 | 483 | 642 |
| FP* | 1428 | 2576 | 2038 | 1544 | 362 | 877 |

* results are for SPEC 2000 benchmarks

Lecture 2
Slide 22

CloudSuite 1.0 (released @ parsa.epfl.ch/cloudsuite)

| | |
|---|---|
| Data Serving Cassandra NoSQL   | MapReduce Machine learning on Hadoop  |
| Media Streaming Apple Quicktime Server   | SAT Solver Symbolic VM constraint  |
| Web Frontend Nginx, PHP server   | Web Search Apache Nutch  |

Lecture 2
Slide 23

Comparing Multiple Programs

| | Computer A | Computer B | Computer C |
|------------------|------------|------------|------------|
| Program 1 (secs) | 1 | 10 | 20 |
| Program 2 (secs) | 1000 | 100 | 20 |
| Program 3 (secs) | 1001 | 110 | 40 |

A is 10 times faster than B for program 1
B is 10 times faster than A for program 2
A is 20 times faster than C for program 1
C is 50 times faster than A for program 2
B is 2 times faster than C for program 1
C is 5 times faster than B for program 2

Each statement above is correct...

...but I just want to know which machine is the best?

Lecture 2
Slide 24

Let's Try a Simpler Example

Two machines timed on two benchmarks

| | Machine A | Machine B |
|-----------|------------|-----------|
| Program 1 | 2 seconds | 4 seconds |
| Program 2 | 12 seconds | 8 seconds |

- How much faster is Machine A than Machine B?
- Attempt 1: **ratio of run times, normalized to Machine A times**

program1: $4/2$ program2: $8/12$

- Machine A ran 2 times faster on program 1, $2/3$ times faster on program 2
- On **average**, Machine A is $(2 + 2/3) / 2 = 4/3$ times faster than Machine B

It turns this “**averaging**” stuff can fool us; watch...

Lecture 2
Slide 25

Example (con't)

Two machines timed on two benchmarks

| | Machine A | Machine B |
|-----------|------------|-----------|
| Program 1 | 2 seconds | 4 seconds |
| Program 2 | 12 seconds | 8 seconds |

- How much faster is Machine A than B?
- Attempt 2: **ratio of run times, normalized to Machine B times**

program 1: $2/4$ program 2 : $12/8$

- Machine A ran program 1 in $1/2$ the time and program 2 in $3/2$ the time
- On **average**, $(1/2 + 3/2) / 2 = 1$
- Put another way, Machine A is **1.0 times faster** than Machine B

Lecture 2
Slide 26

Example (con't)

Two machines timed on two benchmarks

| | Machine A | Machine B |
|-----------|------------|-----------|
| Program 1 | 2 seconds | 4 seconds |
| Program 2 | 12 seconds | 8 seconds |

- How much faster is Machine A than B?
- Attempt 3: **ratio of run times, aggregate (total sum) times, norm. to A**

- Machine A took 14 seconds for both programs
- Machine B took 12 seconds for both programs
- Therefore, Machine A takes $14/12$ of the time of Machine B
- Put another way, Machine A is **6/7 faster** than Machine B

Lecture 2
Slide 27

Which is Right?

Question:

- How can we get **three different answers**?

Solution

- Because, while they are all **reasonable** calculations...

...each answers a different question

We need to be more precise in understanding and posing these performance & metric questions

Lecture 2
Slide 28

Arithmetic and Harmonic Mean

Average of the execution time that tracks total execution time is the **arithmetic mean**

$$\frac{1}{n} \sum_{i=1}^n Time_i$$

This is the defn for "average" you are most familiar with

If performance is expressed as a **rate**, then the average that tracks total execution time is the **harmonic mean**

$$\frac{n}{\sum_{i=1}^n \frac{1}{Rate_i}}$$

This is a different defn for "average" you are prob. less familiar with

Lecture 2
Slide 29

Problems with Arithmetic Mean

Applications do not have the same probability of being run

Longer programs weigh more heavily in the average

For example, two machines timed on two benchmarks

| | Machine A | Machine B |
|-----------|------------------|-----------------|
| Program 1 | 2 seconds (20%) | 4 seconds (20%) |
| Program 2 | 12 seconds (80%) | 8 seconds (80%) |

- If we do arithmetic mean, Program 2 "counts more" than Program 1
 - an improvement in Program 2 changes the average more than a proportional improvement in Program 1
- But perhaps Program 2 is 4 times more likely to run than Program 1

Lecture 2
Slide 30

Weighted Execution Time

Often, one runs some programs more often than others. Therefore, we should **weight** the more frequently used programs' execution time

$$\sum_{i=1}^n Weight_i \times Time_i$$

Weighted Harmonic Mean

$$\frac{1}{\sum_{i=1}^n \frac{Weight_i}{Rate_i}}$$

Lecture 2
Slide 31

Using a Weighted Sum (or weighted average)

| | Machine A | Machine B |
|-----------|------------------|-----------------|
| Program 1 | 2 seconds (20%) | 4 seconds (20%) |
| Program 2 | 12 seconds (80%) | 8 seconds (80%) |
| Total | 10 seconds | 7.2 seconds |

- Allows us to determine relative performance $10/7.2 = 1.38$
--> Machine B is 1.38 times faster than Machine A

Lecture 2
Slide 32

Quiz

E.g., 30 mph for first 10 miles, 90 mph for next 10 miles

What is the average speed?

Average speed = $(30 + 90)/2 = 60$ mph (*Wrong!*)

The correct answer:

Average speed = total distance / total time
 $= 20 / (10/30 + 10/90)$
 $= 45$ mph

For rates use Harmonic Mean!

Lecture 2
Slide 33

Another Solution

Normalize runtime of each program to a reference

| | <u>Machine A (ref)</u> | <u>Machine B</u> |
|-----------|------------------------|------------------|
| Program 1 | 2 seconds | 4 seconds |
| Program 2 | 12 seconds | 8 seconds |
| Total | 10 seconds | 7.2 seconds |

| | <u>Machine A (norm to B)</u> | <u>Machine B (norm to A)</u> |
|-----------|----------------------------------|----------------------------------|
| Program 1 | 0.5 | 2.0 |
| Program 2 | 1.5 | 0.666 |
| Average? | 1.0 | 1.333 |

So when we normalize A to B, and average, it looks like A & B are the same.
 But when we normalize B to A, it looks like A is better!

Lecture 2
Slide 34

Geometric Mean

Used for relative rate or performance numbers

$$Relative_Rate = \frac{Rate}{Rate_{ref}} = \frac{Time_{ref}}{Time}$$

Geometric mean

$$\sqrt[n]{\prod_{i=1}^n Relative_Rate_i} = \frac{\sqrt[n]{\prod_{i=1}^n Rate_i}}{Rate_{ref}}$$

Lecture 2
Slide 35

Using Geometric Mean

| | <u>Machine A (norm to B)</u> | <u>Machine B (norm to A)</u> |
|----------------|----------------------------------|----------------------------------|
| Program 1 | 0.5 | 2.0 |
| Program 2 | 1.5 | 0.666 |
| Geometric Mean | 0.866 | 1.155 |

$1.155 = 1/0.8666!$

Drawbacks:

- Does not predict runtime because it normalizes
- Each application now counts equally

Lecture 2
Slide 36

Well.....

Geometric mean of ratios is not proportional to total time

- ❑ Arithmetic mean in example says machine B is 1.166 times faster
- ❑ Geometric mean says they machine B is 1.155 times faster

Rule of thumb: Use AM for times, HM for rates, GM for ratios

Lecture 2
Slide 37

SPEC Uses Geometric Mean

Steps:

1. for each benchmark i , look up $T_{base,i}$
2. for each benchmark i , run target machine to get $T_{new,i}$

3. compute geometric mean: $\sqrt[n]{\prod_{i=1}^n T_{base,i} / T_{new,i}}$

Is SPEC a good predictor of performance?

Lecture 2
Slide 38

Summary

Performance is important to measure

- ❑ For architects comparing different deep mechanisms
- ❑ For developers of software trying to optimize code, applications
- ❑ For users, trying to decide which machine to use, or to buy

Performance metric are subtle

- ❑ Easy to mess up the “machine A is XXX times faster than machine B” numerical performance comparison
- ❑ You need to know exactly what you are measuring: time, rate, throughput, CPI, cycles, etc
- ❑ You need to know how combining these to give aggregate numbers does different kinds of “distortions” to the individual numbers
- ❑ No metric is perfect, so lots of emphasis on standard benchmarks today

Lecture 2
Slide 39