


Lecture 6 High B/W Caches

Fall 2016

Pejman Lotfi-Kamran



Adapted from slides originally developed by Profs. Falsafi, Hill, Hoe, Lipasti, Shen, Smith, Sohi, and Vijaykumar of Carnegie Mellon University, EPFL, Purdue University, and University of Wisconsin.

Lecture 6
Slide 1

Where Are We?

Fr	Sa	Su	Mo	Tu
	27-Shahrivar		29-Shahrivar	
	3-Mehr		5-Mehr	
	10-Mehr		12-Mehr	
	17-Mehr		19-Mehr	
	24-Mehr		26-Mehr	
	1-Aban		3-Aban	
	8-Aban		10-Aban	
	15-Aban		17-Aban	
	22-Aban		24-Aban	
	29-Aban		1-Azar	
	6-Azar		8-Azar	
	13-Azar		15-Azar	
	20-Azar		22-Azar	
	27-Azar		29-Azar	
	4-Dey		6-Dey	

This Lecture

- High-B/W Caches

Next Lecture:

- Prefetching

Lecture 6
Slide 2

Multi-Level Caches


Processors getting faster w.r.t. main memory

- larger caches to reduce frequency of more costly misses
 - but larger caches are too slow for processor
- => gradually reduce cost of misses with a multiple cache levels

$$t_{avg} = t_{hit} + \text{miss ratio} \times t_{miss}$$

Lecture 6
Slide 3

Multi-Level Cache Design



different requirements
=> different choice of
capacity
block size
associativity

$$t_{avg-L1} = t_{hit-L1} + \text{miss-ratio}_{L1} \times t_{avg-L2}$$

$$t_{avg-L2} = t_{hit-L2} + \text{miss-ratio}_{L2} \times t_{memory}$$

Lecture 6
Slide 4

The Inclusion Property

Multi-level inclusion holds if L2 is always a superset of L1, and so on down the hierarchy. Why?

- if an addr is in L1, then it must be frequently used
- L2 can handle external coherence checks without L1

Causes interesting interactions between L1I, L1D, and L2

Inclusion takes effort to maintain

- each L2 block frame needs status bits marking subblocks in cached by L1
(# of bit = L2 blocksize/L1 blocksize)
- if a L2 block is to be displaced, must first flush out all of its L1-cached subblocks

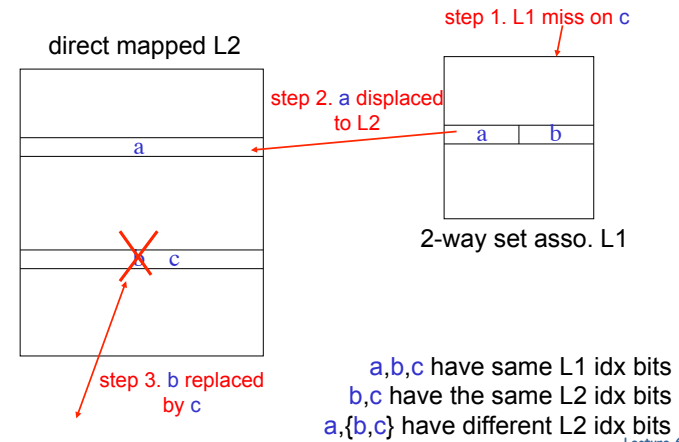
Why would an L2 block be displaced

if it is still has subblocks in L1?

Consider 1. L1 block size < L2 block size 2. associativity in L1

Lecture 6
Slide 5

Possible Inclusion Violation



Multi-Level Inclusion (Cont.)

Inclusion takes effort to maintain

- make L2 cache have bits or pointers giving L1 contents
- invalidate from L1 before replacing from L2
- number of pointers per L2 block
 - L2 blocksize/L1 blocksize

Interesting interaction between L1I, L1D, and L2

- Example a matrix multiply program for a large matrix
- What happens if you maintain inclusion for L1I?

Lecture 6
Slide 7

Level Two Cache Design

What is miss ratio?

- global: L2 misses after L1 references
- local: L2 misses after L1/L1 misses
- solo: as only cache / references

L1 cache design similar to single-level cache

- when main memory “faster” w.r.t. CPU

Apply previous experience to L2 design?

L2 “global” miss ratios not significantly altered by L1 presence

- if L2 cache size $\geq 8 \times$ L1 cache size
- Przybylski et al., ISCA 1989

But L2 caches bigger than before

Lecture 6
Slide 8

Level Two Cache Example

Recall adding associativity to single-level cache helped if

$$\text{diff}(t_{\text{cache}}) + \text{diff}(\text{miss}) \times t_{\text{miss}} < 0$$

$$\text{diff}(\text{miss}) = -1\%, t_{\text{miss}} = 20$$

$$\Rightarrow \text{diff}(t_{\text{cache}}) < 0.2 \text{ cycle}$$

Consider doing the same in an L2 cache where

$$t_{\text{avg}} = t_{\text{cache1}} + \text{miss1} \times t_{\text{cache2}} + \text{miss2} \times t_{\text{memory}}$$

Improvement only if

$$\text{miss1} \times \text{diff}(t_{\text{cache2}}) + \text{diff}(\text{miss2}) \times t_{\text{memory}} < 0$$

$$\text{diff}(t_{\text{cache2}}) < (-\text{diff}(\text{miss2})/\text{miss1}) \times t_{\text{memory}}$$

$$\text{diff}(t_{\text{cache2}}) < 0.0005/0.05 \times 100 = 1 \text{ cycle}$$

Lecture 6
Slide 9

Lock-up Free Caches

Also known as non-blocking caches

Proposed first by Kroft

Only makes sense if processor

- ▢ handle mult. pending references
- ▢ can do useful work under a miss
- ▢ has misses that can be overlapped

Key implementation problems

- ▢ handle reads to pending miss
- ▢ handle writes to pending miss
- ▢ keep multiple requests straight

Memory Access Stream

ld A	hit
ld B	miss
ld C	miss
ld D	hit
st B	miss (pend.)

Miss Status Holding Registers



Lecture 6
Slide 10

Lock-up Free Caches (Cont.)

MSHRs: miss status holding registers

1. Is there already a miss to the same block?
2. Route data back to CPU

- ▢ Valid bit and tag: associatively compared on each miss
- ▢ Status & pointer to block frame
- ▢ What happens on a miss?
- ▢ Tag L1 requests to allow out-of-order service from L2
- ▢ Split-transaction/pipelined L1-L2 interface
- ▢ Associative MSHRs could become bottleneck

Lecture 6
Slide 11

Issue Bandwidth

Increasing issue width => wider caches

Parallel cache access is harder than parallel FUs

- ▢ fundamental difference: caches have state, FUs don't
- ▢ one port affects future for other ports

Several approaches used

- ▢ true multi-porting
- ▢ multiple cache copies
- ▢ virtual multi-porting
- ▢ multi-banking (interleaving)
- ▢ line buffers

Lecture 6
Slide 12

Multi-Port Caches: True Multiporting

Superscalar processors requires multiple data references per cycle

Time-multiplex a single port (double pump)


- need cache access to be faster than datapath clock
- not scalable

Truly multiported SRAMs are

available, but

- more chip area
- slower access

(very undesirable for L1-D)



Lecture 6
Slide 13


Virtual Multi-Porting

Used in Power2 and Alpha 21164

- 21164 uses wave pipelining

Time-share a single port

- may require access to be faster than clock
- probably not scalable beyond 2



Lecture 6
Slide 14

Multiple Cache Copies


Used in Alpha 21264

Independent fast load paths

Single shared store path

Not a scalable solution

- Store is a bottleneck
- Doubles area



Lecture 6
Slide 15


Multi-Banking (Interleaving) Caches

Address space is statically partitioned and assigned to different caches

Which addr bit to use for partitioning?

A compromise (e.g. Intel P6, MIPS R10K)


- multiple references per cyc. if no conflict
 - only one reference goes through if conflicts are detected
 - the rest are deferred
- (bad news for scheduling logic)*



Most helpful is compiler knows about the interleaving rules

Lecture 6
Slide 16

Line Buffers (L0)



Allows multiple ops to/from a block
Increases bandwidth for wide issue

Lecture 6
Slide 17

Juan et al.'s Simulation Study

Focus on two solutions

- true multiporting: ideal
- multi-banking: more scalable

Use dynamically scheduled processor

- 4-issue, 64-inst window to 32 issue, 512-inst window

Direct-mapped 32KB cache

- processor-cache ports
- n-banks
- cell-port (internal)

Lecture 6
Slide 18

Multi-banked Caches

Configuration

- each bank is blocking but interleaving gives non-blocking
- addresses use simple interleaving
- use 4 processor-to-cache ports

Lecture 6
Slide 19

Performance of Multi-banked Caches

Bank conflicts is a problem

- multi-ported vs. multi-banked non-blocking

Bank blocking hurts for small # of banks

- multi-banked non-blocking vs. multi-banked blocking


Lecture 6
Slide 20

Nature of Bank Conflicts

- Many requests are to the same bank
But many are to the same block
- 0.75 of pending requests with 4 banks are to same line

Lecture 6
Slide 21

Bank Conflicts



Lecture 6
Slide 22

Alternative Design

- Multi-banks with multi-ports per bank
- improves performance considerably
 - compare 4 banks 1 port with 2 banks 2 ports
- Add FIFOs to banks
- absorbs some bank busy stalls
- Add coalescing to FIFOs
- combines requests to same line

Lecture 6
Slide 23

Area/Performance

- Multi-bank with same-line optimization seems to work best
- Standard multi-bank caches have relatively low cost
- but also lower performance
- Hybrid designs perform well but at relatively high cost

Lecture 6
Slide 24

Latency vs. Bandwidth

Latency can be handled by

- ▢ hiding/tolerating techniques
 - e.g., parallelism
 - may increase bandwidth demand
- ▢ reducing techniques

Ultimately limited by physics

Lecture 6
Slide 25

Latency vs. Bandwidth (Cont.)

Bandwidth can be handled by

- ▢ banking/interleaving/multiporting
- ▢ wider buses
- ▢ hierarchies (multiple levels)

What happens if average demand not supplied?

- ▢ bursts are smoothed by queues
 - if burst is much larger than average => long queue
 - eventually increases delay to unacceptable levels

Lecture 6
Slide 26