

# Advanced Computer Architecture

## Evaluating Systems

Fall 2016

Pejman Lotfi-Kamran



Adapted from slides originally developed by Profs. Hill, Hoe, Falsafi and Wenisch of CMU, EPFL, Michigan, Wisconsin

Fall 2016

Lec.11 - Slide 1

## Where Are We?

Fr	Sa	Su	Mo	Tu
	27-Shahrivar		29-Shahrivar	
	3-Mehr		5-Mehr	
	10-Mehr		12-Mehr	
	17-Mehr		19-Mehr	
	24-Mehr		26-Mehr	
	1-Aban		3-Aban	
	8-Aban		10-Aban	
	15-Aban		17-Aban	
	22-Aban		24-Aban	
	29-Aban		1-Azar	
	6-Azar		8-Azar	
	13-Azar		15-Azar	
	20-Azar		22-Azar	
	27-Azar		29-Azar	
	4-Dey		6-Dey	

◆ This Lecture  
● Evaluation

◆ Next Lecture:  
● Evaluation

Fall 2016

Lec.11 - Slide 2

## Roadmap

➔ Metrics

◆ Tools

◆ Workloads

Fall 2016

Lec.11 - Slide 3

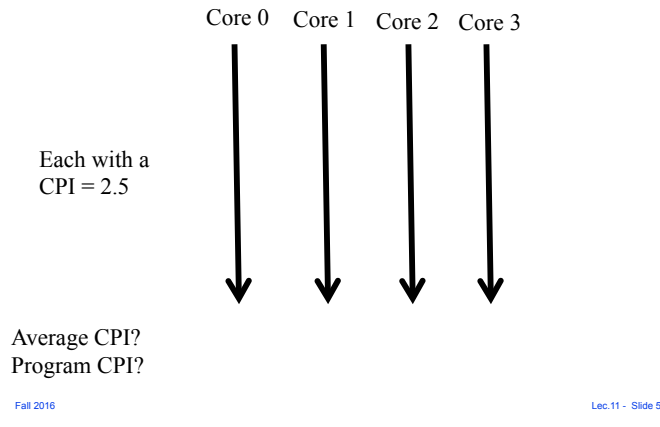
## Performance Metric

- ◆ **Uniprocessor**: Cycles per instruction (CPI)
- ◆ **Multiprocessor**: CPI is not proportional to perf.
  - Threads often synchronize through spinning
  - Waiting for others to catch up
  - Or allowing access to shared data one at a time
- ◆ Spinning threads have low CPI!!!!
  - Averaging across threads would lead to misleading results
- ◆ In the limit, if all threads are waiting all the time
  - Average CPI is really low
  - But, there is no forward progress (the program is running sequentially)
  - Not all instructions make forward progress
  - Frequent spins on I/O and locks

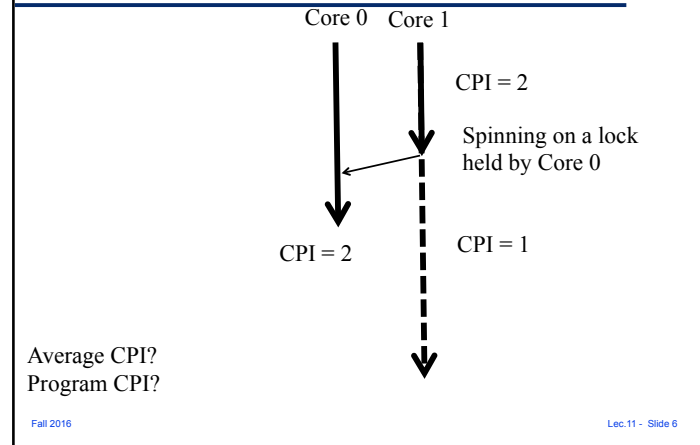
Fall 2016

Lec.11 - Slide 4

### Example: CPI not useful for Parallel Programs



### Example: CPI not useful for Parallel Programs




### CPI not Useful in General

- ◆ Bottomline is execution time
  - ◆ If spinning can be counted out, then CPI works
  - ◆ Transactional server workloads
  - ◆ Mostly spin in the OS
  - ◆ User-level CPI is valid
  - ◆ Can be averaged
- Fall 2016 Lec.11 - Slide 7

### MP performance metric


- ◆ Typical metric: Transactions completed / min. (TPM)
- ◆ Coarse progress metrics bad for sampling
  - Transaction completions infrequent
  - High variance of inter-arrival and service times



*Measuring TPM requires large measurements*

## Fine-grain performance metric

- ◆ User-mode inst. per trans. constant per app. cfg.
- ◆ Hence, user-mode  $IPC \propto TPM$ 
  - Validated for TPC-C, SpecWEB [Hankins 2003, Wenisch 2006]
  - Most apps yield rather than spin in user mode



*Impact: Same confidence with 1000x shorter measurements*

Fall 2016

Lec.11 - Slide 9

## Other Metrics: Cost

- ◆ Cost is a function of more than just the processor.
- ◆ Cost is a complex function of many hardware components and software
- ◆ Cost is often not a "smooth" function
  - Often a function of packaging
    - ▲ how many pins on a board
    - ▲ how many processors on a board
    - ▲ how many boards in a chassis
  - Typically refers to design cost
- ◆ There is also
  - Manufacturing cost
  - Operation cost (becoming huge)

Fall 2016

Lec.11 - Slide 10

## Operation Cost: Enterprise IT (Server) Energy Usage

Kenneth Brill (Uptime Institute)

- ◆ "Economic Meltdown of Moore's Law"
- ◆ In 2012: Energy/server lifetime 50% more than price/server
  - And 2% of all Carbon footprint in the US

Energy Star report to Congress:

- ◆ Datacenter energy 2x from 2000 to 2006
- ◆ Roughly 2% of all electricity & growing

TCO (Total Cost of Operation) is the metric!

Fall 2016

Lec.11 - Slide 11

## Other Metrics: Reliability

Until today

- ◆ mostly dealt with at the circuit/technology level
- ◆ mostly a specialized market
  - if high reliability is important (e.g., ATM transactions)
    - ▲ Deal with it in HW
    - ▲ Large redundant mainframes for banks
    - ▲ Multiple redundant/heterogeneous computers in planes
  - Otherwise tolerate faults and/or handle in software
- ◆ in the future, technologies will be error-prone
  - May have to deal with reliability in all computing market segments
    - ▲ Embedded, handheld, desktop, server, supercomputer
    - ▲ Specialize hardware to protect only parts needed
- ◆ Related to cost (design + energy)

Fall 2016

Lec.11 - Slide 12

## Roadmap

- ◆ Metrics
- ➔ Tools
- ◆ Workloads

Fall 2016

Lec.11 - Slide 13

## How do we evaluate?


- ◆ Model it analytically (analytic models)
  - Fast
  - Not accurate
- ◆ Model it in software (simulator)
  - REAL SLOW
  - More accurate
- ◆ Model it in FPGA
  - Fast
  - Limited in model size
  - Slow to program
- ◆ Hardware counters
  - Fast
  - Not flexible!

Fall 2016

Lec.11 - Slide 14

## Simulation Speed Challenges

- ◆ Longer benchmarks
  - SPEC 2006: Trillions of instructions per benchmark
- ◆ Slower simulators
  - Full-system simulation: 1000× slower than SimpleScalar
- Multiprocessor systems
  - CMP: 2x cores every processor generation



**1,000,000× slowdown vs. HW → years per experiment**

Fall 2016

Lec.11 - Slide 15

## Full-system simulation is slow

- ◆ Simulation slowdown per cpu

● Real HW:	~ 2 GIPS	1 s
● Simics:	~ 30 MIPS	66 s
● Flexus, no timing:	~ 900 KIPS	37 m
● Flexus, OoO:	~ 24 KIPS	23 h

**2 years to simulate 10 seconds of a 64-core workload!**

Fall 2016

Lec.11 - Slide 16

## Our solution: Statistical sampling

### ◆ Measure uniform or random locations



### ◆ Impact

- Sampling: **~10,000×** reduction in turnaround time
- Independent measurements: **100- to 1000-way** parallelism
- Confidence intervals: **quantifiable** result reliability

### ◆ Challenges

- Rapidly create warm  $\mu$ arch state prior to measurements
- Allow independent simulation of each measurement
- Sample non-deterministic, highly variable MP applications

Fall 2016

Lec.11 - Slide 17

## Functional Simulation

### ◆ Functional simulation is faster than detailed simulation

- Flexus (no timing) is 38 times faster than Flexus (OoO)

### ◆ Use functional simulation for “warmup”

- Memory (guarantees correctness)
- Cache hierarchy (avoids bias)
- Branch predictor (avoids bias)



**No state for core microarchitecture → Bias**

Fall 2016

Lec.11 - Slide 18

## Handling Bias

### ◆ Core micro-architecture can be warmed up rapidly

- Detailed simulation to warmup core micro-architecture

### ◆ Perform warmup prior to measurement

- Functional warming during fast-forwarding
- Detailed warmup before each simulation window



Fall 2016

Lec.11 - Slide 19

## Simulation Speedup

### ◆ 10 seconds of a 64-core workload

- Normal execution: 2 years
- With sampling: 20 days



### ◆ 37x improvement in speed but not enough


### ◆ Solution

- Avoid functional simulation (which lasts 17 days)
- Accelerate detailed simulation (which lasts 3 days)

Fall 2016

Lec.11 - Slide 20

## Avoiding Functional Simulation



- ◆ Store warm cache & branch predictor state
  - Same sample design, accuracy, confidence
  - No warming length prediction needed


*Works for any microarchitecture*

Fall 2016

Lec.11 - Slide 21

## Accelerating Detailed Simulation

- ◆ Checkpoint library makes measurement independent
- ◆ Run multiple measurements in parallel



Fall 2016

Lec.11 - Slide 22

## Simulation Speedup

- ◆ Sampling without a checkpoint library:

- 10 seconds of a 64-core workload: 20 days



- ◆ Sampling with a checkpoint library:

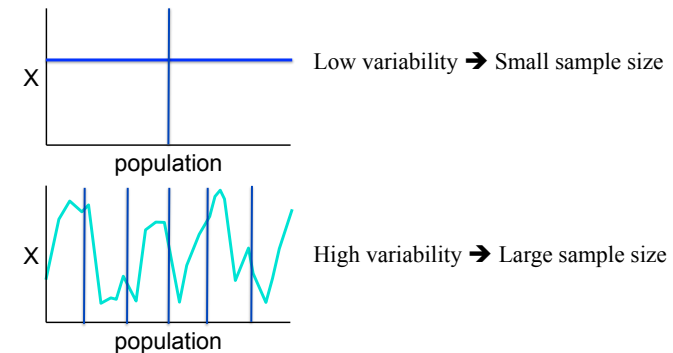
- 10 seconds of a 64-core workload: 3 hours with 100 cores



Fall 2016

Lec.11 - Slide 23

## How to Choose the Sample Size?



*Variability determines sample size*

Fall 2016

Lec.11 - Slide 24

## Steps for Timing Simulation

### 1. Prepare workload for simulation

- Port workload into Simics

### 2. Measure baseline variance

- Determine required library size



### 3. Collect checkpoints

- Via functional warmup



### 4. Detailed Simulation

- Estimate performance results



Fall 2016

Lec.11 - Slide 25

## 2. Determine Sampling Parameters

- ◆ Guess variability
- ◆ Generate flexpoints for the variability
- ◆ Run timing simulation
- ◆ Measure error and correct the guess

Fall 2016

Lec.11 - Slide 26

## Typical Sampling Parameters

	Flexus (64-CPU CMP.OoO)
Warming	100k cycles
Measurement	50k cycles
Target confidence	95%
Sample size	800
Sim. time per checkpoint	~ 20 min


Fall 2016

Lec.11 - Slide 27

## 3. Checkpoint Creation


### ◆ Spread Simics checkpoints

- Simics fast mode rapidly covers 10 seconds



### ▲ Collect flexpoints in parallel

- Via CMP.L2Shared.Trace
- From each Simics checkpoint



Fall 2016

Lec.11 - Slide 28

## 4. Detailed Simulation

- ◆ Run detailed simulation with OoO simulators
- ◆ Process all flexpoints, aggregate offline
- ◆ Manipulate results with *stat-manager*
  - Each run creates binary `stats_db.out` database
  - Offline tools to select subsets; aggregate
  - Generate text reports from simple templates
  - Compute confidence intervals for mean estimates

Fall 2016

Lec.11 - Slide 29

## Matched-pair comparison [Ekman 05]


- ◆ Often interested in relative performance
- ◆ Change in performance across designs varies less than absolute change
- ◆ Matched pair comparison
  - Allows smaller sample size
  - Reports confidence in performance change

Fall 2016

Lec.11 - Slide 30

## Matched-pair example

Performance results for two microarchitecture designs  
checkpoints processed in random order




**Lower variability in performance delta reduces sample size by 3.5 to 150x**

Fall 2016

Lec.11 - Slide 31

## Matched-pair with Flexus

- ◆ Simple  $\mu$ Arch changes (e.g., changing latencies)
  - use same flex-points
- ◆ Complex changes (e.g., adding components)



Fall 2016

Lec.11 - Slide 32