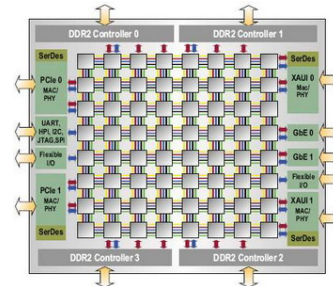


Advanced Computer Architecture

Advanced Coherence Fall 2016

Pejman Lotfi-Kamran



Adapted from slides originally developed by Profs. Hill, Hoe, Falsafi and Wenisch of CMU, EPFL, Michigan, Wisconsin

Fall 2016

Lec.15 - Slide 1

Where Are We?

Fr	Sa	Su	Mo	Tu
	27-Shahrivar		29-Shahrivar	
	3-Mehr		5-Mehr	
	10-Mehr		12-Mehr	
	17-Mehr		19-Mehr	
	24-Mehr		26-Mehr	
	1-Aban		3-Aban	
	8-Aban		10-Aban	
	15-Aban		17-Aban	
	22-Aban		24-Aban	
	29-Aban		1-Azar	
	6-Azar		8-Azar	
	13-Azar		15-Azar	
	20-Azar		22-Azar	
	27-Azar		29-Azar	
	4-Dey		6-Dey	

◆ This Lecture

- Advanced Coherence
 - ▲ Efficient snooping
 - ▲ Directories
 - ▲ Optimization

◆ Next Lecture:

- Consistency

Fall 2016

Lec.15 - Slide 2

A Note on L2 Inclusion

◆ If L2 inclusive

- On a WB, update copy in L2
- Upon subsequent reads, data comes from L2

◆ If L2 non-inclusive

- A read might find other L1 copies
- Can get data from L1
- Or go to the next level

Fall 2016

Lec.15 - Slide 3

No snoop, but centralize around directory

◆ No snoop because of the network is not a point of serialization

◆ But, make directory the point of serialization

- ▲ Requests for block A are serialized at directory
- ▲ Requests across blocks (e.g., for block B) can overlap

◆ For example

- Read miss for A blocks until read completes
- Read miss for A from another CPU blocks while the first is pending
- Read miss for B proceeds while A is pending

Fall 2016

Lec.15 - Slide 4

CMP Directories

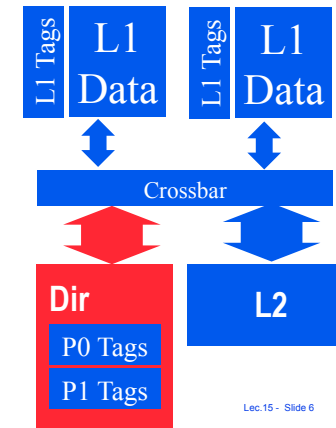
- ◆ Duplicate Tag
 - Keep a copy of all of L1's Tag Arrays
- ◆ Sparse
 - Keep a limited number of Tags
 - FULL: Per-tag keep bit vector
 - Partial: Keep a limited list of sharers

Fall 2016

Lec.15 - Slide 5

Duplicate Directories

- ◆ Keeps a copy of all L1 Tags
 - Compaq Piranha, ARM Elbea Chip (ARM A9)
- ◆ Upon request (GetR, GetX, GetUpgr)
 - Searches all L1 Tags
 - if found, MSI transitions
 - If not, go to L2

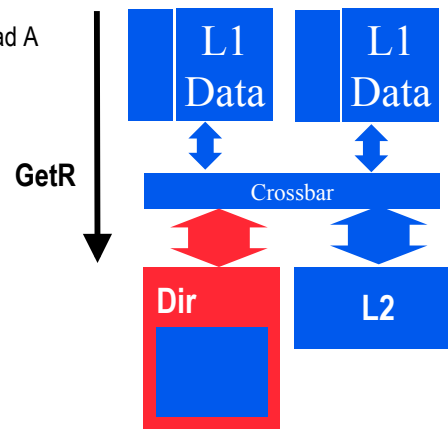


Fall 2016

Lec.15 - Slide 6

Example

- ◆ P0 Read A

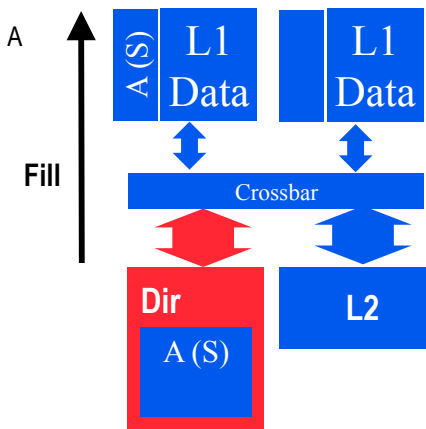


Fall 2016

Lec.15 - Slide 7

Example

- ◆ P0 Read A

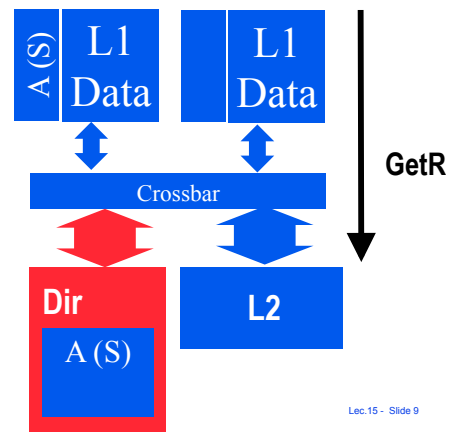


Fall 2016

Lec.15 - Slide 8

Example

- ◆ P0 Read A
- ◆ P1 Read A

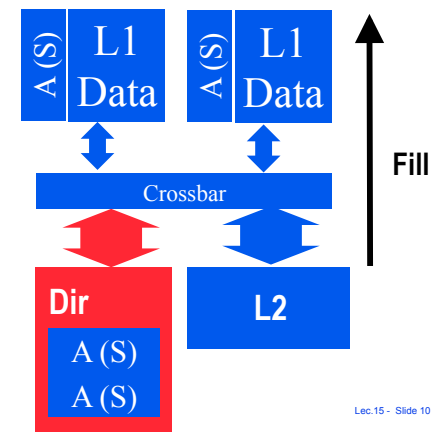


Fall 2016

Lec.15 - Slide 9

Example

- ◆ P0 Read A
- ◆ P1 Read A

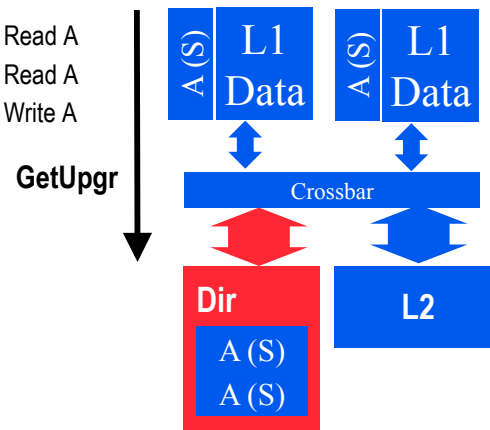


Fall 2016

Lec.15 - Slide 10

Example

- ◆ P0 Read A
- ◆ P1 Read A
- ◆ P0 Write A

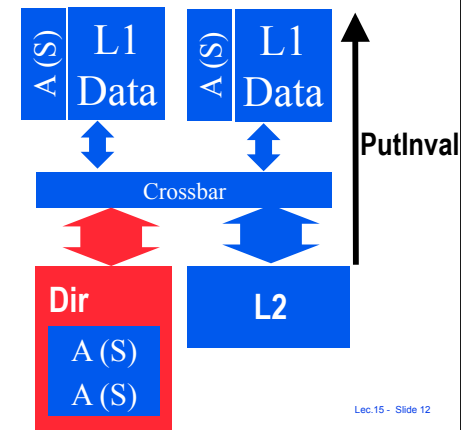


Fall 2016

Lec.15 - Slide 11

Example

- ◆ P0 Read A
- ◆ P1 Read A
- ◆ P0 Write A

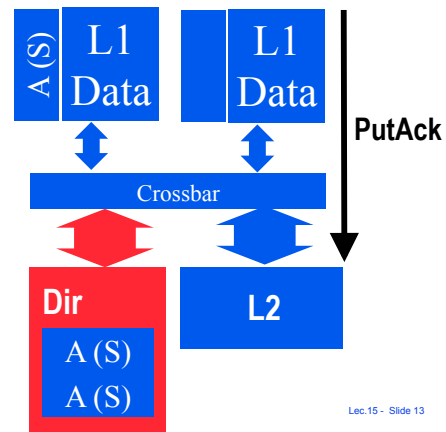


Fall 2016

Lec.15 - Slide 12

Example

- ◆ P0 Read A
- ◆ P1 Read A
- ◆ P0 Write A

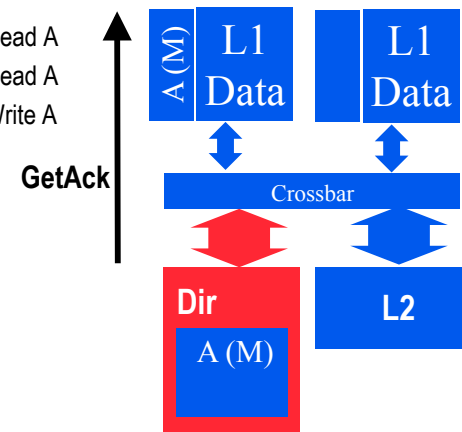


Fall 2016

Lec.15 - Slide 13

Example

- ◆ P0 Read A
- ◆ P1 Read A
- ◆ P0 Write A



Fall 2016

Lec.15 - Slide 14

A Note on Deadlock

- ◆ Transactions are:
 - Non-atomic and interleaved
- ◆ In the case of a general-purpose NoC
 - To avoid deadlock
 - Use request/response virtual channels
 - As long as requests are serviced while waiting, deadlock can be avoided

Fall 2016

Lec.15 - Slide 15

Duplicate Tags Problem (1)

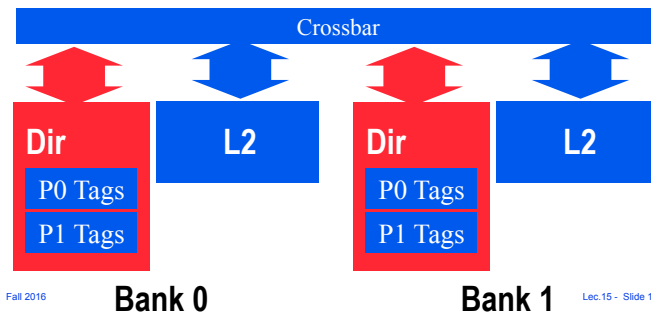
- ◆ Centralized
 - A centralized structure is not much better than the bus
 - Directory becomes a bottleneck
 - But, large L2's are banked (multi-banked, address-interleaved)
 - Solution: distribute directories with L2 banks

Fall 2016

Lec.15 - Slide 16

Distributed Directories

- ◆ Bank based on block address
- ◆ Example: Odd address (1), even addresses (2)



Distributed directories

- ◆ Distribute the traffic across banks
- ◆ But suffers from the same problems
 - Bank conflict if burst of accesses to one bank
- ◆ As with caches, in practice works well
 - E.g., Compaq Piranha chip, 8 L2 + directory banks

Fall 2016

Lec.15 - Slide 18

Duplicate Tags Problem (2)

- ◆ Don't scale!
- ◆ Assume 4-way L1's
- ◆ Each directory search is N times 4 ways
- ◆ For two L1's (two cores) → 8 ways to search
- ◆ For 32 L1's (32 cores) → 128 ways to search
 - Bottleneck in latency
 - Bottleneck in energy

Fall 2016

Lec.15 - Slide 19

Sparse Directories

- ◆ Do not keep all tags
- ◆ Keep a cache of tags
- ◆ Limited/fixed associativity
 - Invalidate blocks in L1 which you can't track in directory
- ◆ Represent duplicate tags with a bit vector
 - Copies of one tag have bits set in vector

Fall 2016

Lec.15 - Slide 20

Example: Duplicate Tag

- ◆ P0 has cached blocks A (shared) and B (modified)
- ◆ P1 has cached blocks A (shared) and C (shared)
- ◆ A, B and C all fall in the same set in L1's

Duplicate Tag

P0,A(S)	P0,B(M)	P1,C(S)	P1,A(S)
---------	---------	---------	---------

- ◆ All entries for one L1 set map to the same set in directory
- ◆ Search associatively

Fall 2016

Lec.15 - Slide 21

Example: Sparse (3-way associative)

P0 Read A

Sparse Tag (3-way)

A(S)[01]		
----------	--	--

Fall 2016

Lec.15 - Slide 22

Example: Sparse (3-way associative)

P0 Read A
P1 Read C

Sparse Tag (3-way)

A(S)[01]	C(S)[10]	
----------	----------	--

Fall 2016

Lec.15 - Slide 23

Example: Sparse (3-way associative)

P0 Read A
P1 Read C
P1 Read A

Sparse Tag (3-way)

A(S)[11]	C(S)[10]	
----------	----------	--

Fall 2016

Lec.15 - Slide 24

Example: Sparse (3-way associative)

P0 Read A
P1 Read C
P1 Read A
P0 Write B

Sparse Tag (3-way)

A(S)[11]	C(S)[10]	B(M)[01]
----------	----------	----------

Fall 2016

Lec.15 - Slide 25

Example: Sparse (2-way associative)

P0 Read A

Sparse Tag (2-way)

A(S)[01]	
----------	--

Fall 2016

Lec.15 - Slide 26

Example: Sparse (2-way associative)

P0 Read A
P1 Read C

Sparse Tag (2-way)

A(S)[01]	C(S)[10]
----------	----------

Fall 2016

Lec.15 - Slide 27

Example: Sparse (2-way associative)

P0 Read A
P1 Read C
P1 Read A

Sparse Tag (2-way)

C(S)[10]	A(S)[11]
----------	----------

Fall 2016

Lec.15 - Slide 28

Example: Sparse (2-way associative)

P0 Read A
P1 Read C
P1 Read A
P0 Write B (invalidate C from P1)

Sparse Tag (2-way)

B(M)[01]	A(S)[11]
----------	----------

Fall 2016

Lec.15 - Slide 29

Summary

- ◆ Buses don't scale
 - Use directories
 - Point-to-point messages
 - Serialize per address at directory
- ◆ Centralized directories are a bottleneck
 - Distribute
- ◆ Can't keep all tags (duplicates)
 - Sparse directories

Fall 2016

Lec.15 - Slide 30

Roadmap

- ◆ Snoopy optimization
- ◆ Directory protocols
- ◆ Protocol optimization
- **Directory optimization**

Fall 2016

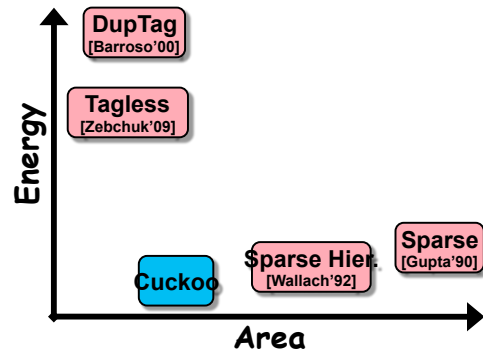
Lec.15 - Slide 31

Scaling Coherence Directories

- ◆ CMPs heading to high core counts
 - High performance requires private caches
 - Private caches require coherence
 - Coherence limited by directory conflicts
- ◆ Traditional ways to reduce conflicts
 - Increase associativity (Dup-Tag) → pay Power
 - Over-provision sets (Sparse) → pay Area

Want low-power and area-efficient organization

Directory Efficiency



Cuckoo Directory: low power, small area

Cuckoo Directory Organization

- ◆ Inspired by Cuckoo Hash [Pagh'01, Fotakis'03]
- ◆ Re-locates conflicting entries
 - Cuckoos put eggs in others' nests
 - Move "native" eggs to make space
- ◆ Eliminates practically all conflicts
 - Low associativity (3-way)
 - Small area



Fall 2016

Lec.15 - Slide 34

Contributions

- ◆ Leverage Cuckoo Hash organization
 - Algorithm to re-locate conflicting entries
 - 3-way, but behaves like fully-associative
- ◆ Cuckoo Directory
 - More power-efficient than Dup-Tag
 - More area-efficient than Sparse
 - Scalable to 1000+ cores

Cuckoo Directory: a truly scalable directory

Outline

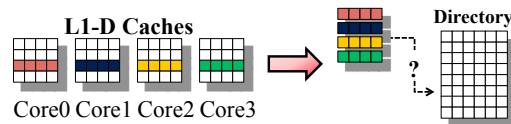
- ◆ Introduction
- ◆ Directory Conflicts and Solutions
- ◆ Cuckoo Directory
- ◆ Evaluation
- ◆ Conclusions

Fall 2016

Lec.15 - Slide 36

Origin of the Conflict Problem

- ◆ Directory maps all blocks in private caches
- ◆ Low-order bits in “same” set identical
- ◆ Entries map to same directory indexes
 - either... increase associativity with each core
 - or... over-provision to spread across sets

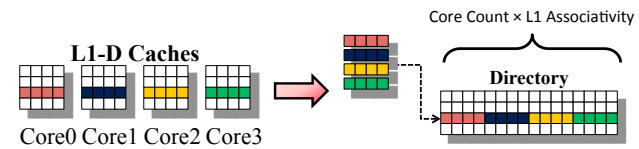


Fall 2016

Lec.15 - Slide 37

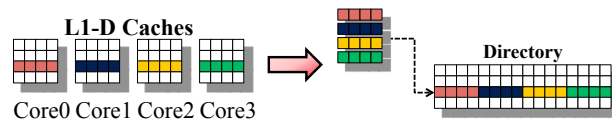
Duplicate-Tag Directory

- ◆ Ensure sufficient space for worst-case
 - Allocate space for each way of each cache
- ◆ High core counts → huge associativity



Sparse Directory

- ◆ Split L1 sets into directory sets
 - Grow number of bits in index

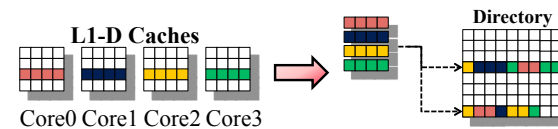


Fall 2016

Lec.15 - Slide 39

Sparse Directory

- ◆ Split L1 sets into directory sets
 - Grow number of bits in index

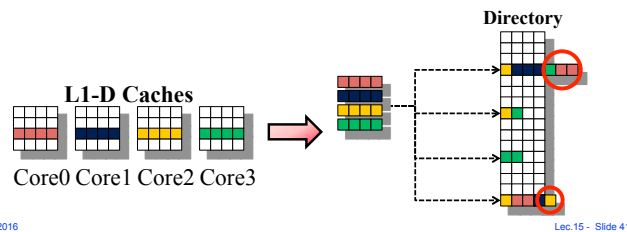


Fall 2016

Lec.15 - Slide 40

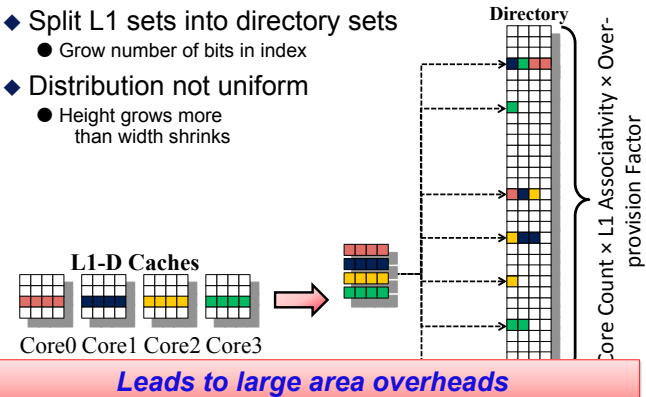
Sparse Directory

- ◆ Split L1 sets into directory sets
 - Grow number of bits in index
- ◆ Distribution not uniform



Sparse Directory

- ◆ Split L1 sets into directory sets
 - Grow number of bits in index
- ◆ Distribution not uniform
 - Height grows more than width shrinks



Methodology

- ▶ **Flexus** simulation infrastructure [Wenisch'06]
- ▶ Full-system trace (OS + user) evaluation

Benchmark Applications

- OLTP: TPC-C
 - IBM DB2 & Oracle
- DSS: TPC-H
 - IBM DB2 Qrys 2, 16, 17
- Web: SPECweb99
 - Apache & Zeus
- Scientific
 - em3d & ocean

Model Parameters

- L1 (I & D) : 64KB, 2-way
- L2: 16MB (1MB per core)

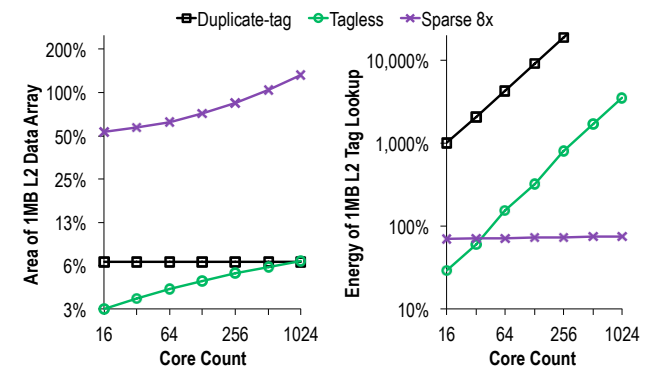
Area Model

- Total bit count

Energy Model

- Total bit accessed
- Reads and Writes equal

Area and Energy Comparison



No design is both area- and power-efficient

Outline

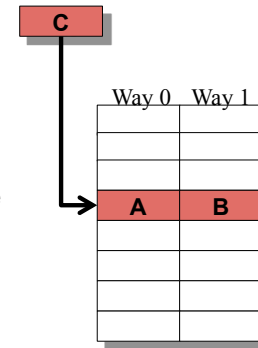
- ◆ Introduction
- ◆ Directory Conflicts and Solutions
- ◆ Cuckoo Directory
- ◆ Evaluation
- ◆ Conclusions

Fall 2016

Lec.15 - Slide 45

Sparse Directory

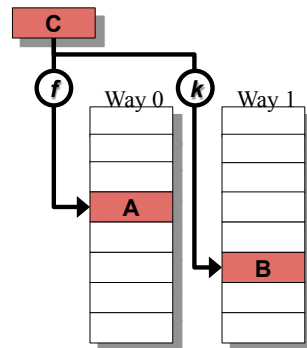
- ◆ C conflicts with A
- ◆ C conflicts with B
- ◆ Therefore...
 - A, B, C can't coexist
 - Even if space available



To insert C, directory must evict A or B

Cuckoo Directory

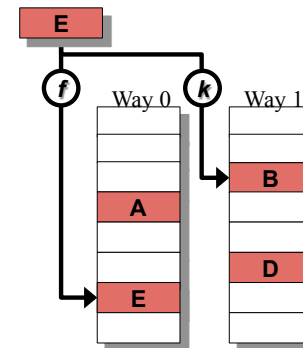
- ◆ Skewed structure [Seznec'93]
 - Hash each way's index
- ◆ **KEY IDEA**
 - Move conflicting entries
 - Inserting C moves A



A, B, and C can coexist

Example 1: Common Case

- ◆ Case
 - $f(E)$ or $k(E)$ available
- ◆ Action
 - Insert immediately



Fall 2016

Lec.15 - Slide 48

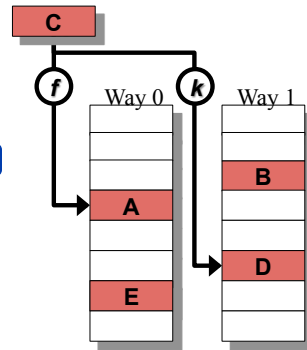
Example 2: Uncommon Case

◆ Case

- $f(C)$ and $k(C)$ occupied

◆ Action

- Place C, replacing A
- Place A, replacing B
- Place B



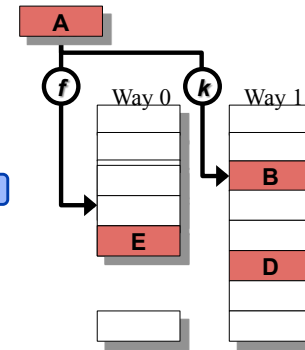
Example 2: Uncommon Case

◆ Case

- $f(C)$ and $k(C)$ occupied

◆ Action

- Place C, replacing A
- Place A, replacing B
- Place B



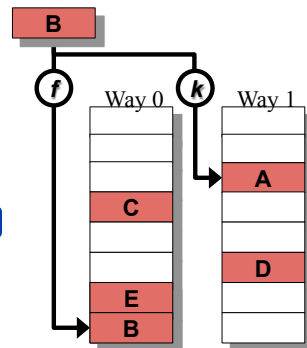
Example 2: Uncommon Case

◆ Case

- $f(C)$ and $k(C)$ occupied

◆ Action

- Place C, replacing A
- Place A, replacing B
- Place B

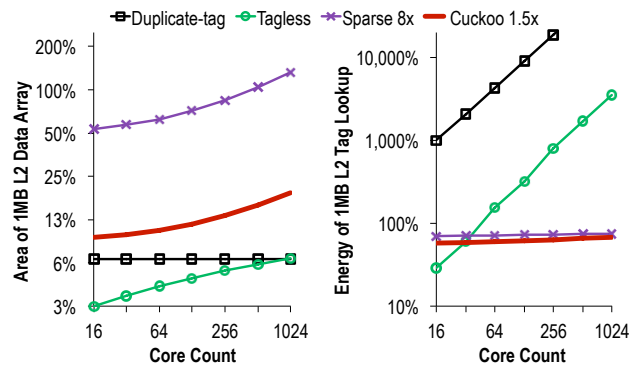


All entries fit in directory without eviction

Outline

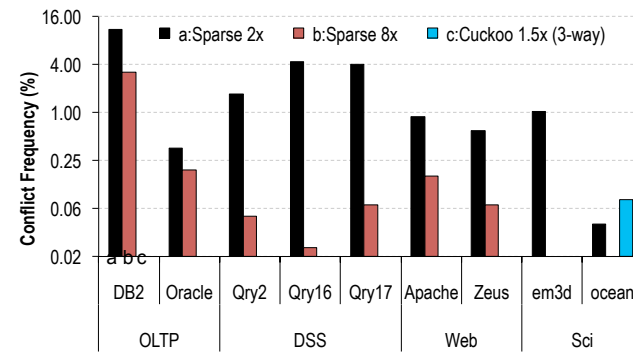
- ◆ Introduction
- ◆ Directory Conflicts and Solutions
- ◆ Cuckoo Directory
- ◆ Evaluation
- ◆ Conclusions

Area and Energy Comparison



Cuckoo Directory is area- AND power-efficient

Directory Conflict Frequency



0.08% conflicts in worst-case benchmark

Conclusions

- ◆ Cuckoo Directory
 - Moves conflicting entries across ways
 - More power-efficient than Dup-Tag & Tagless
 - More area-efficient than Sparse & Coarse
- ◆ Scalable to 1000+ cores
 - Constant associativity for any core count
 - Minimal capacity over-provisioning

Cuckoo Directory: a truly scalable directory