

# Advanced Computer Architecture

## Memory Ordering

Fall 2016

Pejman Lotfi-Kamran



Adapted from slides originally developed by Profs. Hill, Hoe, Falsafi and Wenisch of CMU, EPFL, Michigan, Wisconsin

Fall 2016

Lec.17 - Slide 1

## Where Are We?

Fr	Sa	Su	Mo	Tu
	27-Shahrivar		29-Shahrivar	
	3-Mehr		5-Mehr	
	10-Mehr		12-Mehr	
	17-Mehr		19-Mehr	
	24-Mehr		26-Mehr	
	1-Aban		3-Aban	
	8-Aban		10-Aban	
	15-Aban		17-Aban	
	22-Aban		24-Aban	
	29-Aban		1-Azar	
	6-Azar		8-Azar	
	13-Azar		15-Azar	
	20-Azar		22-Azar	
	27-Azar		29-Azar	
	4-Dey		6-Dey	

◆ This Lecture  
● Consistency (2)

◆ Next Lecture:  
● Consistency (3)

Fall 2016

Lec.17 - Slide 2

## The Big Debate

When shall hardware enforce order?

**Always order:**

- Always (SC)
  - ▲ Software wants a “multiprogrammed CPU” behavior
- Easy to program & understand
- Assumed to exhibit inferior performance

**Relax order:**

- Order enforced through software annotation
- Let the hardware overlap write latency

Fall 2016

Lec.17 - Slide 3

## Big Misconception: Large Performance Gap

But, strong ordering thought to hurt performance

- One reason for a variety of memory models and flavors

Not true!

Memory only has to **appear** to be ordered

- Hardware can relax order speculatively
- Save state while speculating
- Roll back if relaxed order observed by others
- E.g. result, SC + Speculation  $\geq$  RC!

This is the Bart Simpson’s approach to relaxing order:  
“I didn’t do it. Noone saw me doing it!”

Fall 2016

Lec.17 - Slide 4

## Evolution of SC Systems

Naive SC: every access is ordered

Optimized SC:

- Three simple optimizations
- Existing pipeline HW
- E.g., MIPS R10K

Wait-free SC:

- SC++
- ASO
- InvisiFence

Fall 2016

Lec.17 - Slide 5

## Enhancing SC's Performance

1. Review: Store buffering
  - Store misses (either missing data or permission) can be removed from pipeline in program order
  - Place them in a store buffer
2. Review: All miss "fetches" can be overlapped
  - L1 cache is a point of serialization (think "switch")
  - Requests for cache blocks can be sent out in parallel
  - All accesses to L1 must be atomic and in order
  - Therefore:
    - Order in which blocks are fetched/filled has no bearing on actual observed program order
    - Not until L1 cache access for the block is performed
3. New: Load hits can "speculate" past store misses
  - MIPS R10000 did this!

Fall 2016

Lec.17 - Slide 6


## Review: SC + Store Buffer + Fetch blocks

ST X	; removed from pipeline (wait in SB)
ST B	; removed from pipeline (wait in SB)
LD A	; looks up cache but waits (no load)
ALU	; waiting for LD A
ST A	; looks up cache but waits (no store)
LD Y	; looks up cache, fetches Y into L1
ST Z	; instruction not fetched

Fall 2016

Lec.17 - Slide 7

## SC + Store Buffer + Fetch + Speculative Load



Loads A speculatively [Gharachorloo '91, Ranganathan '97]

+ LD A, ALU and ST A are done

**Not enough buffering (ROB full)!**

Fall 2016

Lec.17 - Slide 8

## MIPS R10K: SC + Store Buffer + Fetch + Speculative Load

**ST X** ; removed from pipeline (wait in SB)  
**ST B** ; removed from pipeline (wait in SB)  
**LD A** ; speculatively loaded  
**ALU** ; Done!  
**ST A** ; waiting for store buffer  
**LD Y** ; looks up cache, fetches Y into L1  
**ST Z** ; instruction not fetched

Fall 2016

Lec.17 - Slide 9

## What if another processor asks for A?

LD A is out of order with respect to ST X & ST B!

What if another processor stores A meanwhile?

- Coherence message sent to this processor
- LD/ST queue has a speculative bit for LD A
- If the bit is set → **misspeculation!**


On misspeculation:

- Service the coherence message (e.g., invalidate A)
- Re-execute starting from the LD A
- LD A will turn into a miss
- Noone saw the reordering (except for the local processor)!

Fall 2016

Lec.17 - Slide 10

## Misspeculation in MIPS R10K



Store to A from another processor

- ◆ Flush all ROB and LD/ST queue entries
- ◆ Re-execute from LD A

Fall 2016

Lec.17 - Slide 11

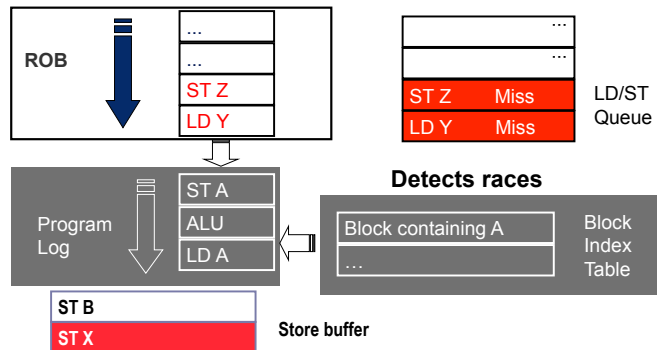
## Requirements for Store-Stall-Free SC: SC++ [Gniady'99]

- ① Non-blocking caches
- ② Logging program state while a store is pending
  - May need to log thousands of cycles
- ③ Fast lookup to detect order violation
  - Upon request for data by others
- ④ Infrequent rollbacks
  - Typical of well-behaved parallel applications
  - Rollbacks are due to false sharing or data races

Fall 2016


Lec.17 - Slide 12

## SC++: A Store-Stall-Free SC



- Maintain a “log” of computation, with a fast block lookup
- All order relaxed up to first missing LD (LD Y)!

## Performance Comparison



- Data from RSIM (left) and Flexus (right)
  - 16 1-GHz OoO processors, Piranha DSM
  - TSO = Stores can not bypass each other
- Up to 70% gap between SC & SC++
- SC++ > RC >> TSO

Slide 14

## Can we do any model ++?

- ◆ Relaxed models conservatively enforce order at Fences!
  - If there are may fences in the code performance hurts
  - E.g., XCHG instructions in PC (TSO)
  - E.g., Fence instructions in WC & RC
- ◆ PC++, WC++, RC++?
  - Yes!
  - Can use the same buffering and lookup
  - Relax Fence instructions!

Fall 2016

Lec.17 - Slide 15

## Summary

- ◆ Can relax order as long as
  - There is enough buffering
  - The buffer can be looked up for races (coherence messages)
- ◆ Relaxed models may not be necessary to achieve high performance!
  - Hardware can relax order
  - Models may be needed to allow compilers reorder
  - Compilers have not had much impact on performance

Fall 2016

Lec.17 - Slide 16