

Advanced Computer Architecture

Memory Ordering

Fall 2016

Pejman Lotfi-Kamran



Adapted from slides originally developed by Profs. Hill, Hoe, Falsafi and Wenisch of CMU, EPFL, Michigan, Wisconsin

Fall 2016

Lec.18 - Slide 1

Where Are We?

Fr	Sa	Su	Mo	Tu
	27-Shahrivar		29-Shahrivar	
	3-Mehr		5-Mehr	
	10-Mehr		12-Mehr	
	17-Mehr		19-Mehr	
	24-Mehr		26-Mehr	
	1-Aban		3-Aban	
	8-Aban		10-Aban	
	15-Aban		17-Aban	
	22-Aban		24-Aban	
	29-Aban		1-Azar	
	6-Azar		8-Azar	
	13-Azar		15-Azar	
	20-Azar		22-Azar	
	27-Azar		29-Azar	
	4-Dey		6-Dey	

Fall 2016

Lec.18 - Slide 2

◆ This Lecture
● Consistency (3)

◆ Next Lecture:
● Synchronization

Atomic sequence ordering (ASO) [Wenisch' 07]

- ◆ Intuition: If access sequence appears atomic, then actual order does not matter
- ◆ Dynamically group accesses into **atomic sequences**
- ◆ Relax order within & across sequences
- ◆ Stores: get all permissions; then commit atomically
 - On-chip commit (into L2)
- ◆ On a race, recover to sequence's checkpoint
 - Races are rare in typical parallel apps. [Gniady' 99]


Coarse-grain rollback enables practical HW

Fall 2016

Lec.18 - Slide 3

Atomic sequence ordering (ASO)

- ◆ Enforce order over coarse-grain atomic sequences



- ◆ Detect races:
 - Mark speculative loads in cache
 - Coherence request to speculative block ⇒ violation
- ◆ CPU rollback: 2-4 CPU checkpoints

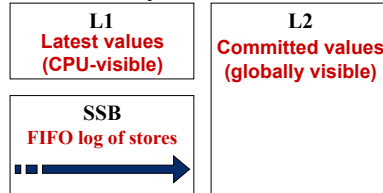
Coarse-grain rollback enables practical HW

Fall 2016

Lec.18 - Slide 4

Scalable Store Buffer (SSB)

- ◆ Maintain consistency at L2 instead of L1




- ◆ Fast access to latest value in L1
- ◆ Commit consistent values from SSB → L2

SSB is FIFO – no RMW, no associative search

Fall 2016

Lec.18 - Slide 5

Operation under ASO



Accumulate accesses into sequence

Await permission for all stores


Commit writes atomically

Obtain permissions in advance → on-chip commit

Fall 2016

Lec.18 - Slide 6

Handling stores: Commit



Store to L1 and append to SSB


Drain updates from SSB to L2

Delay external requests while draining for atomicity

Fall 2016

Lec.18 - Slide 7

Handling stores: Rollback



Detect race on external request

Discard all speculative blocks from L1


Replay stores SSB → L1 (similar to [Gandhi 05])

Forward progress: 1 non-spec. store if 1st seq. rolls back

Fall 2016

Lec.18 - Slide 8

ASO Hardware



- ◆ Fast access to CPU-visible and globally-visible values
 - Stores: issue to L1&SSB; commit SSB→L2
- ◆ Detect races: L1 & L2 track sequences' reads
 - Recover to checkpoint; reconstruct L1 from SSB
- ◆ Similar HW for TLS, TM → add generic mechanisms

Fall 2016

Lec.18 - Slide 9

Evaluation methodology

- ◆ Flexus OoO timing simulation
- ◆ SMARTS statistical sampling

Benchmark Applications

OLTP: TPC-C
 IBM DB2 & Oracle
 DSS: TPC-H Qry 1, 6, 16
 IBM DB2
 Web: SPECweb99
 Apache & Zeus
 Scientific
 barnes, ocean


Model Parameters

16-node directory-based DSM
 4GHz CPUs; SPARC ISA
 4-wide OoO; 8-stage pipeline
 64KB L1, 8MB L2
 Store prefetch & speculative loads [Gharachorloo' 91]

Fall 2016

Lec.18 - Slide 10

Performance comparison



- ◆ Up to 50% gap between aggressive SC and RMO
- ◆ Even under RMO, fences cause 5-10% unnecessary stalls

ASO & SSB enable ≥RMO perf. with SC guarantees

Fall 2016

Lec.18 - Slide 11

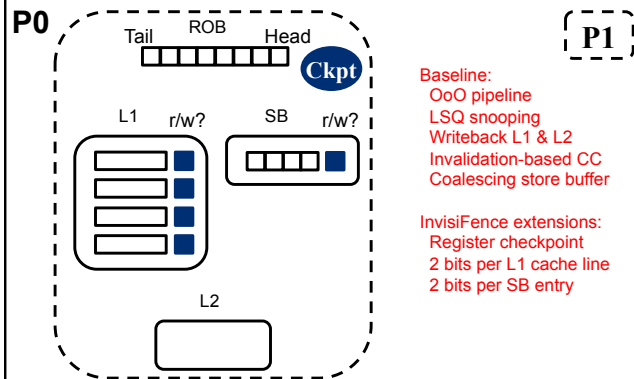
InvisiFence For Weak Ordering

- ◆ Add deep speculation to eliminate stalling on fences
- ◆ Mechanism: register ckpt + 2 bits per L1 cache line
 - Similar HW to other deep speculation (TLS, TM, Cherry...)
- ◆ Initiate speculation at fence instructions
 - Detect violations via cache coherence protocol
 - Preserve non-speculative data in L2 (facilitates rollback)
- ◆ Speculation ends when store buffer becomes empty
 - Commit by flash-clearing read/write bits

F

InvisiFence - Blundell - ISCA 2009

InvisiFence Hardware

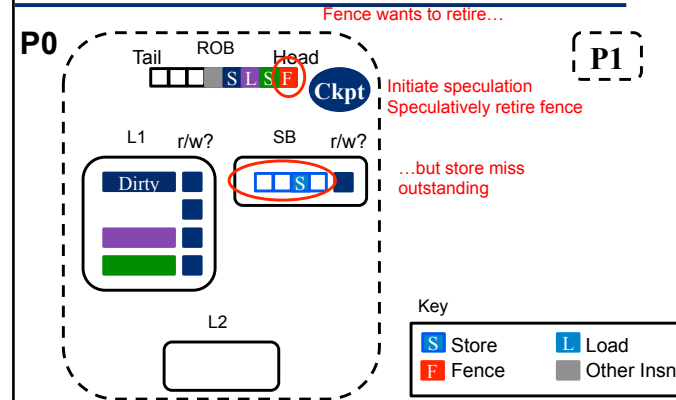


Fall 2016

InvisiFence - Blundell - ISCA 2009

Lec.18 - Slide 13

InvisiFence: Example

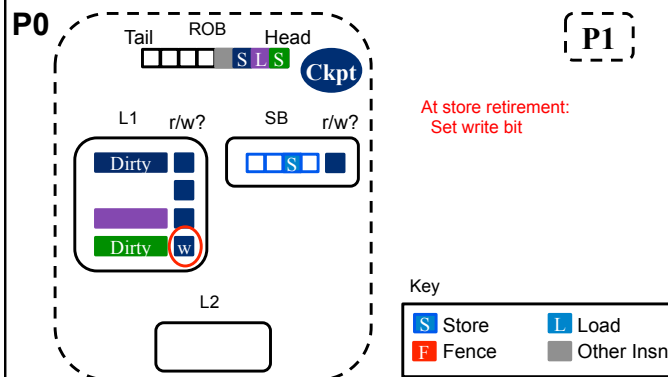


Fall 2016

InvisiFence - Blundell - ISCA 2009

Lec.18 - Slide 14

InvisiFence: Violation Detection

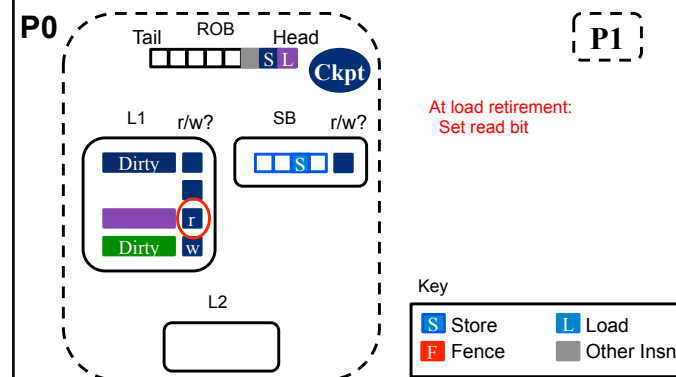


Fall 2016

InvisiFence - Blundell - ISCA 2009

Lec.18 - Slide 15

InvisiFence: Violation Detection

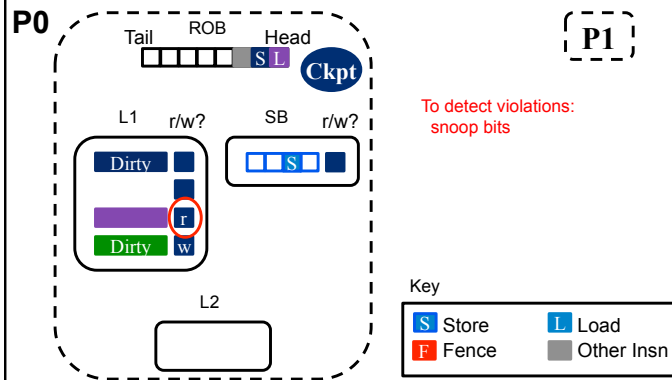


Fall 2016

InvisiFence - Blundell - ISCA 2009

Lec.18 - Slide 16

InvisiFence: Violation Detection

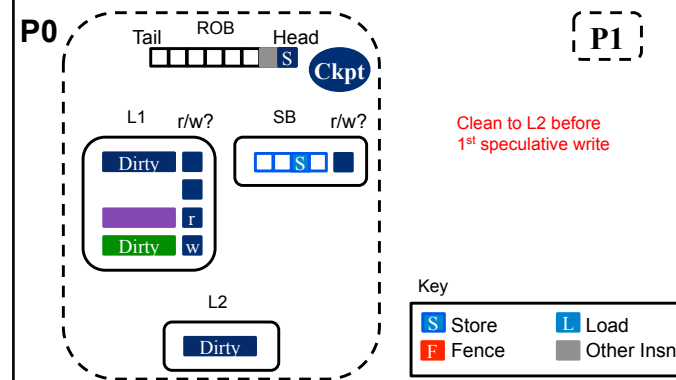


Fall 2016

InvisiFence - Blundell - ISCA 2009

Lec.18 - Slide 17

InvisiFence: Version Management

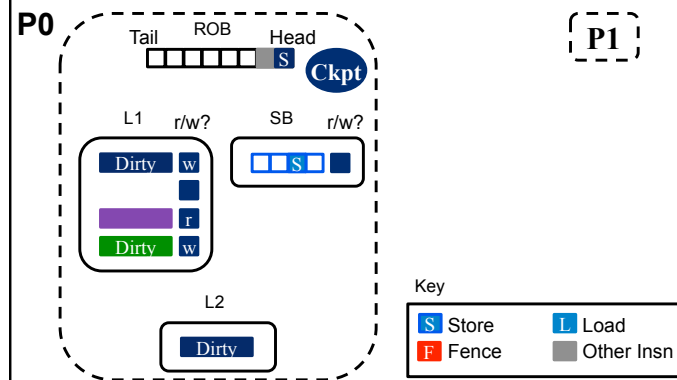


Fall 2016

InvisiFence - Blundell - ISCA 2009

Lec.18 - Slide 18

InvisiFence: Version Management

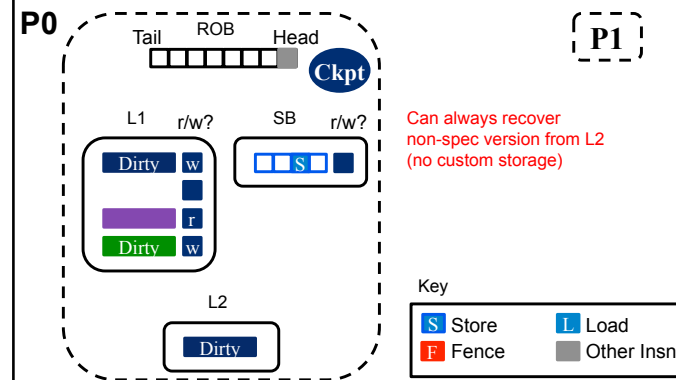


Fall 2016

InvisiFence - Blundell - ISCA 2009

Lec.18 - Slide 19

InvisiFence: Version Management

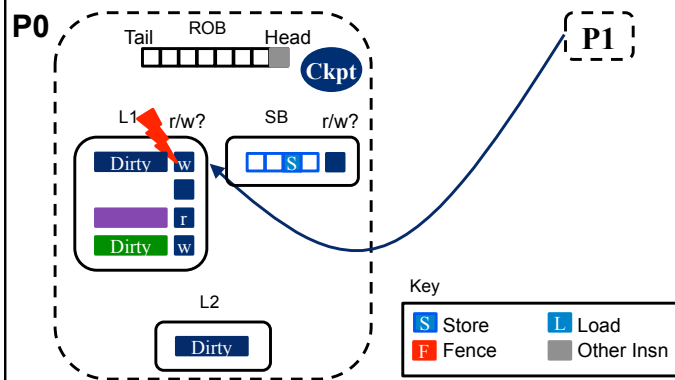


Fall 2016

InvisiFence - Blundell - ISCA 2009

Lec.18 - Slide 20

InvisiFence: Rollback

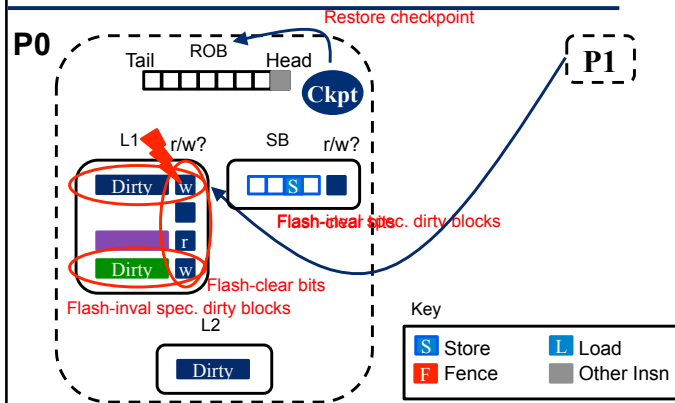


Fall 2016

InvisiFence - Blundell - ISCA 2009

Lec.18 - Slide 21

InvisiFence: Rollback

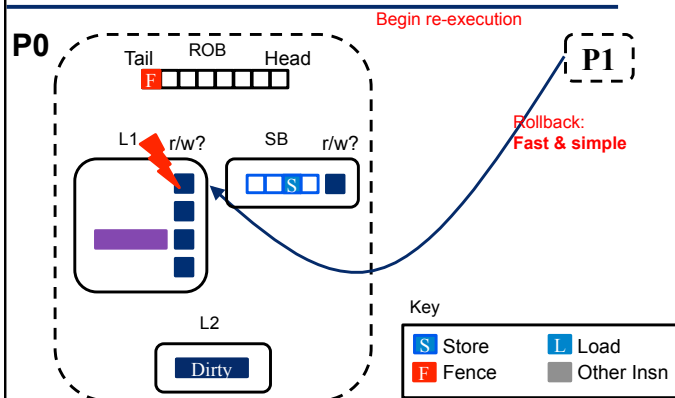


Fall 2016

InvisiFence - Blundell - ISCA 2009

Lec.18 - Slide 22

InvisiFence: Rollback

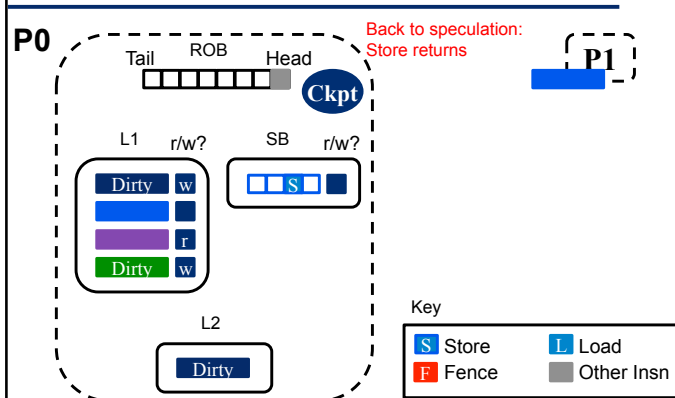


Fall 2016

InvisiFence - Blundell - ISCA 2009

Lec.18 - Slide 23

InvisiFence: When to Commit?




Fall 2016


InvisiFence - Blundell - ISCA 2009

Lec.18 - Slide 24

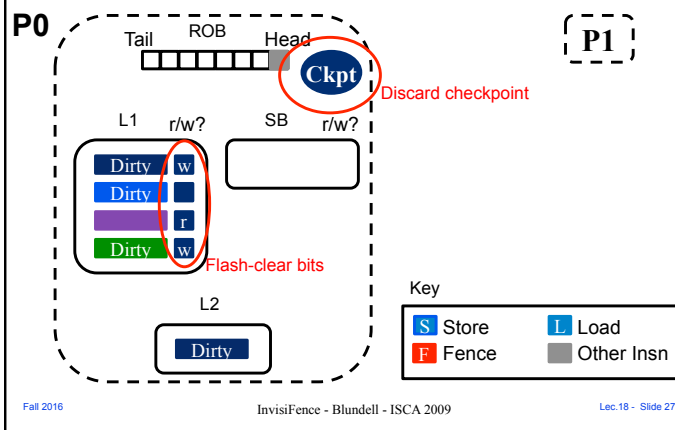
InvisiFence: When to Commit?



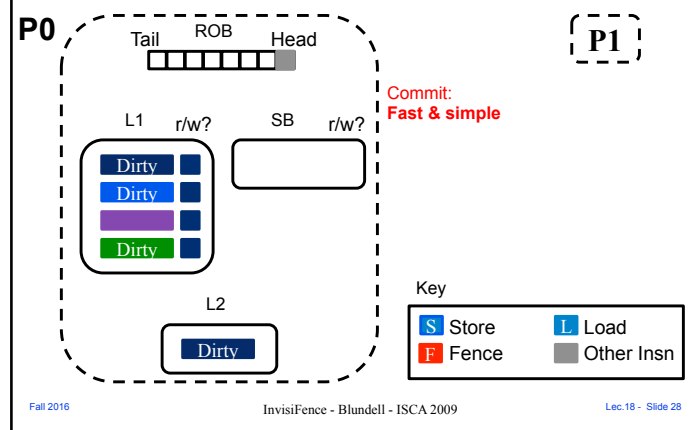
InvisiFence: When to Commit?





InvisiFence: Commit



InvisiFence: Commit






Generalizing InvisiFence for Strong Ordering


- ◆ Strong models impose additional ordering constraints
 - Processor Consistency (x86, TSO): ordering between stores
 - Sequential Consistency: ordering between all operations
- ◆ These constraints are conceptually “implicit fences”
 - e.g., for SC: every operation is “implicit fence”
- ◆ InvisiFence can handle these just like explicit fences!
 - Increases speculation frequency...

No other hardware changes

Fall 2016 InvisiFence - Blundell - ISCA 2009 Lec.18 - Slide 31



Strong Ordering Performance (SC)



Comparison to Atomic Sequence Ordering [Wenisch'07]:
Both eliminate stalls

Fall 2016

InvisiFence - Blundell - ISCA 2009

Lec.18 - Slide 33

ASO & InvisiFence: Design Comparison

ASO [Wenisch'07]

- ◆ Fine-grained tracking
 - 1K-entry store buffer
 - 10 KB
- ◆ Lengthy commit
 - Atomically drain SB to L2
 - Multiple checkpoints
- ◆ Changes to L1
 - Mult. per-block R/W bits
 - Write-through
 - Per-word valid bits

InvisiFence

- ◆ Coalesced tracking
 - 8-entry store buffer
 - < 1 KB
- ◆ Constant-time commit
 - Flash-clear bits
 - Single checkpoint
- ◆ Changes to L1
 - Single per-block R/W bits
 - Clean to L2

Both eliminate stalls, but InvisiFence hardware simpler

Fall 2016

InvisiFence - Blundell - ISCA 2009

Lec.18 - Slide 34

Summary: Speculative Memory Ordering

Memory Consistency models are neither **necessary** nor **sufficient** to achieve high performance!

- ◆ No need for “static” compile-time fences
- ◆ Detect and enforce fences dynamically
- ◆ Speculation techniques effective
- ◆ Similar mechanisms needed for Transactional Memory or Lock speculation

Fall 2016

Lec.18 - Slide 35