

# Advanced Computer Architecture

## Synchronization

Fall 2016

Pejman Lotfi-Kamran



Adapted from slides originally developed by Profs. Hill, Hoe, Falsafi and Wenisch of CMU, EPFL, Michigan, Wisconsin

Fall 2016

Lec.20 - Slide 1

## Where Are We?

Fr	Sa	Su	Mo	Tu
	27-Shahrivar		29-Shahrivar	
	3-Mehr		5-Mehr	
	10-Mehr		12-Mehr	
	17-Mehr		19-Mehr	
	24-Mehr		26-Mehr	
	1-Aban		3-Aban	
	8-Aban		10-Aban	
	15-Aban		17-Aban	
	22-Aban		24-Aban	
	29-Aban		1-Azar	
	6-Azar		8-Azar	
	13-Azar		15-Azar	
	20-Azar		22-Azar	
	27-Azar		29-Azar	
	4-Dey		6-Dey	

- ◆ This Lecture
  - Synchronization

- ◆ Next Lecture:
  - TM

Fall 2016

Lec.20 - Slide 2

## Lock Elision

Slides are courtesy Dan Sorin of Duke University

Fall 2016

Lec.20 - Slide 3

## Speculative Lock Elision

- ◆ Multithreaded Programming gains Importance
  - SMPs, Multithreaded Architectures
  - Used in non-traditional fields (Desktop)
- ◆ Synchronization Mechanisms required for exclusive access
  - Serialize access to shared data structures
  - Necessary to prevent conflicting updates

Fall 2016

Lec.20 - Slide 4

## Dynamically Unnecessary Synchronization

```
LOCK (locks->error_lock)
if (local_error > multi-
    >err_multi)
    multi-
    >err_multi = local_err;
UNLOCK (locks-
    >error_lock);
```

Thread 1

```
LOCK (hash_tbl.lock)
var = hash_tbl.lookup (X)
if (!var)
    hash_tbl.add (X);
UNLOCK (hash_tbl.lock)
```

Thread 2

```
LOCK (hash_tbl.lock)
var = hash_tbl.lookup (Y)
if (!var)
    hash_tbl.add (Y);
UNLOCK (hash_tbl.lock)
```

- ◆ Require no lock, if
  - no write access
  - different fields accessed

Fall 2016

Lec 20 - Slide 5

## Multithreaded Programming

- ◆ Conservative Locking
  - Easier to show correctness
  - Lock more often than necessary
- ◆ Locking Granularity
  - Trade-Off between Performance and Complexity
- ◆ Thread-unsafe legacy libraries
  - Require global locking

Fall 2016

Lec 20 - Slide 6

## False Dependencies

- ◆ Locks introduce Control and Data Dependence
- ◆ Removal
  - Key is *Appearance of Instantaneous Change*
  - Lock can be elided, if memory operation remain atomic
    - ▲ Data read is not modified by other threads
    - ▲ Data written is not access by other threads
  - Any instruction that violates these conditions must not be retired

Fall 2016

Lec 20 - Slide 7

## Silent store pairs

	Program Semantic	Instruction Stream	Value of _lock_	
◆ i16 & i6 are silent pair			as seen by self	as seen by other threads
● i16 <i>undoes</i> i6				
◆ Not silent individually	TEST _lock_	L1: i1 ld1 t0, 0(t1)	FREE	FREE
● No write to _lock_		i2 bne t0, L1:		
● Read _lock_ is ok	TEST _lock_	i3 ld1_l t0, 0(t1)	FREE	FREE
	&	i4 bne t0, L1:		
	SET _lock_	i5 lda t0, 1(0)		
◆ SLE does not depend on semantic info		i6 stl_c t0, 0(t1)	HELD	FREE
● Simply observe silent pairs		i7 beq t0, L1:		
		i8- i15		
	<i>critical section</i>			
	RELEASE _lock_	i16 stl 0, 0(t1)	FREE	FREE

(C) 2008 Babak Falsafi from Adve, Falsafi, Hill, Lebeck, Reinhardt, Smith & Singh

## Two SLE Predictions

### ◆ Silent pair prediction

- On a store
  - ▲ Predict another store will undo changes
  - ▲ Don't perform stores
  - ▲ Monitor memory location
- If validated, elide two stores

### ◆ Atomicity prediction

- Predict all memory access within silent pair occur atomically
- No partial updates visible to other threads

Fall 2016

Lec.20 - Slide 9

## Initiating Speculation

### ◆ Detect candidate pairs

- Filter indexed by program counter
- Add confidence metric for each pair

### ◆ Predict, if lock held

- If yes, don't speculate

Fall 2016

Lec.20 - Slide 10

## Buffering Speculative State

### ◆ Register state

- Reorder Buffer
  - ▲ Already used for branch prediction
- Register Checkpoint

### ◆ Memory state

- Augment write buffers
  - ▲ Keep speculative writes in write buffer...
  - ▲ ...until validated
  - ▲ Allows merging of speculative writes

Fall 2016

Lec.20 - Slide 11

## Misspeculation conditions

### ◆ Atomicity violation

- Detected by coherence protocol
  - ▲ Sufficient for ROB register state
  - ▲ For register checkpoint add speculative access bit to L1 cache

### ◆ Resource constraints

- Cache Size
- Write Buffer Size

Fall 2016

Lec.20 - Slide 12

## Committing Speculative Memory State

- ◆ Commits must appear instantaneously
- ◆ Cache state and Cache data
  - Can speculate on state
  - Can't speculate on data
- ◆ Speculative store
  - Send GETX request
  - Block already exclusive when speculative writes commit


Fall 2016

Lec 20 - Slide 13

## Example (from Lang Yiu)

```
LOCK(lock->error_lock)
if (local_err > multi->err_multi)
    multi->err_multi = local_err;
UNLOCK(lock->error_lock)
```

```
L1: 1. ldl t0, 0(t1)      #t0 = lock
    2. bne t0, L1:        #if not free, goto L1
    3. ldl_l t0, 0(t1)    #load locked, t0 = lock
    4. bne t0, L1:        #if not free, goto L1
    5. lda t0, 1(0)       #t0 = 1
    6. stl_c t0, 0(t1)    #conditional store, lock = 1
    7. beq t0, L1:        #if stl_c failed, goto L1
    8. ldq t0, 0(s4)
    9. ldc $f10, 0(t0)
    10. cmplt $f10, $f11, $f10
    11. fbeq $f10, L2:
    12. stt $f11, 0(t0)
    13. ldq t1, -31744(gp)
    14. ldq t0, 0(t1)
    15. ldq t1, 32(t0)
    16. stl 0, 0(t1)      #lock = 0, release lock
```




Fall 2016

Lec 20 - Slide 14

## How to elide locks? - Example

```
LOCK(lock->error_lock)
if (local_err > multi->err_multi)
    multi->err_multi = local_err;
UNLOCK(lock->error_lock)

L1: 1. ldl t0, 0(t1)      #t0 = lock
    2. bne t0, L1:        #if not free, goto L1
    3. ldl_l t0, 0(t1)    #load locked, t0 = lock
    4. bne t0, L1:        #if not free, goto L1
    5. lda t0, 1(0)       #t0 = 1
    6. stl_c t0, 0(t1)    #conditional store, lock = 1
    7. beq t0, L1:        #if stl_c failed, goto L1
    8. ldq t0, 0(s4)
    9. ldc $f10, 0(t0)
    10. cmplt $f10, $f11, $f10
    11. fbeq $f10, L2:
    12. stt $f11, 0(t0)
    13. ldq t1, -31744(gp)
    14. ldq t0, 0(t1)
    15. ldq t1, 32(t0)
    16. stl 0, 0(t1)      #lock = 0, release lock
```




Fall 2016

Lec 20 - Slide 15

## How to elide locks? - Example

```
LOCK(lock->error_lock)
if (local_err > multi->err_multi)
    multi->err_multi = local_err;
UNLOCK(lock->error_lock)

L1: 1. ldl t0, 0(t1)      #t0 = lock
    2. bne t0, L1:        #if not free, goto L1
    3. ldl_l t0, 0(t1)    #load locked, t0 = lock
    4. bne t0, L1:        #if not free, goto L1
    5. lda t0, 1(0)       #t0 = 1
    6. stl_c t0, 0(t1)    #conditional store, lock = 1
    7. beq t0, L1:        #if stl_c failed, goto L1
    8. ldq t0, 0(s4)
    9. ldc $f10, 0(t0)
    10. cmplt $f10, $f11, $f10
    11. fbeq $f10, L2:
    12. stt $f11, 0(t0)
    13. ldq t1, -31744(gp)
    14. ldq t0, 0(t1)
    15. ldq t1, 32(t0)
    16. stl 0, 0(t1)      #lock = 0, release lock
```



Fall 2016

Lec 20 - Slide 16


## How to elide locks? - Example

```

LOCK(lock->error_lock)
if (local_err > multi->err_multi)
    multi->err_multi = local_err;
UNLOCK(lock->error_lock)

L1: ldl t0, 0(t1)      #t0 = lock
    bne t0, L1         #if not free, goto L1
    ldl t1, 0(t1)      #load locked, t0 = lock
    mrrw t0, L1        #if not free, goto L1
    lds t0, 1(t0)       #t0 = 1
    stl_c t0, 0(t1)     #conditional store, lock = 1
    beq t0, L1         #if stl_c failed, goto L1
    ldg t0, 0(s4)
    ldc $f10, 0(t0)
    cmplt $f10, $f11, $f10
    fbeq $f10, L2
    stt $f11, 0(t0)
    ldg t1, -31744(gp)
    ldg t0, 0(t1)
    ldg t1, 32(t0)
    stl 0, 0(t1)        #lock = 0, release lock

```



Fall 2016

Lec 20 - Slide 17

## Evaluation

### ◆ 3 Systems

- CMP, SMP and DSM
- Total Store Order
- Single Register Checkpoint
- 32 entry lock predictor

### ◆ Run on SimpleMP

- Based on Simplescalar


### ◆ Benchmarks

- SPLASH
  - ▲ mp3d, barnes, cholesky
- SPLASH-2
  - ▲ Radiosity, water-nsq, ocean
- Microbenchmark

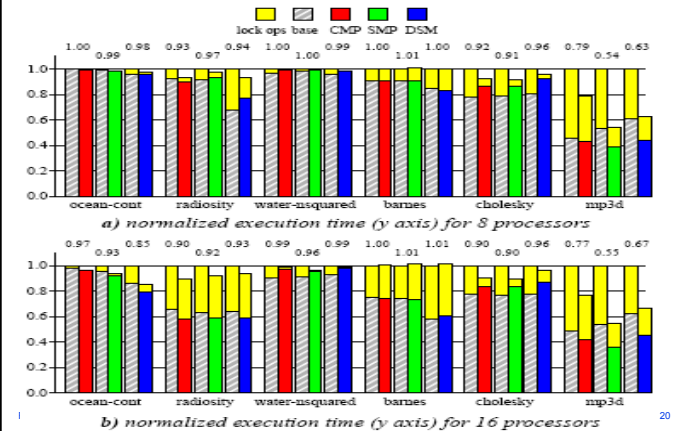
Fall 2016

Lec 20 - Slide 18

## Lock acquires/releases elided



## SPLASH Performance



## Asymmetric Critical Section Execution

Slides from M. Aater Suleman et. al., ASPLOS 2009

Fall 2016

Lec.20 - Slide 21

## Background

- ◆ To leverage CMPs:
  - Programs must be split into *threads*
- ◆ Mutual Exclusion:
  - Threads are not allowed to update shared data concurrently
- ◆ Accesses to shared data are encapsulated inside **critical sections**
- ◆ Only one thread can execute a critical section at a given time

Fall 2016

Lec.20 - Slide 22

## Example of Critical Section from MySQL

List of Open Tables


	A	B	C	D	E
Thread 0	X			X	
Thread 1			X		
Thread 2	X	X			
Thread 3				X	X

Thread 2:  
CloseAllTables()

Fall 2016

Lec.20 - Slide 23

## Example Critical Section from MySQL



Fall 2016

Lec.20 - Slide 24

## Example Critical Section from MySQL

LOCK\_open → Acquire()


End of Transaction:  
foreach (table opened by thread)  
if (table.temporary)  
table.close()

LOCK\_open → Release()

Fall 2016

Lec 20 - Slide 25

## Contention for Critical Sections




Fall 2016

Lec 20 - Slide 26

## Impact of Critical Sections on Scalability

- Contention for critical sections increases with the number of threads and limits scalability



MySQL (oltp-1)

Fall 2016

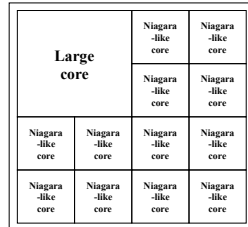
## Outline

- ◆ Background
- ◆ **Mechanism**
- ◆ Performance Trade-Offs
- ◆ Evaluation
- ◆ Related Work and Summary

Fall 2016

Lec 20 - Slide 28

## The Asymmetric Chip Multiprocessor (ACMP)



- Provide one large core and many small cores
- Execute parallel part on small cores for high throughput
- Accelerate serial part using the large core

Fall 2016


Lec 20 - Slide 29

## Conventional ACMP

```

EnterCS()
PriorityQ.insert(...)
LeaveCS()
    
```

1. P2 encounters a Critical Section
2. Sends a request for the lock
3. Acquires the lock
4. Executes Critical Section
5. Releases the lock




Fall 2016

Lec 20 - Slide 30

## Accelerating Critical Sections (ACS)

**Critical Section  
Request Buffer  
(CSRB)**



Accelerate Amdahl's serial part & **critical sections** using the large core

Fall 2016


Lec 20 - Slide 31

## Accelerated Critical Sections (ACS)

```

EnterCS()
PriorityQ.insert(...)
LeaveCS()
    
```

1. P2 encounters a Critical Section
2. P2 sends CSCALL Request to CSRB
3. P1 executes Critical Section
4. P1 sends CSDONE signal



Fall 2016

Lec 20 - Slide 32



## Architecture Overview


- ◆ ISA extensions
  - CSCALL LOCK\_ADDR, TARGET\_PC
  - CSRET LOCK\_ADDR
- ◆ Compiler/Library inserts CSCALL/CSRET
- ◆ On a CSCALL, the small core:
  - Sends a CSCALL request to the large core
    - ▲ Arguments: Lock address, Target PC, Stack Pointer, Core ID
  - Stalls and waits for CSDONE
- ◆ Large Core
  - Critical Section Request Buffer (CSRB)
  - Executes the critical section and sends CSDONE to the requesting core

Fall 2016

Lec.20 - Slide 33

## False Serialization

- ◆ ACS can serialize independent critical sections
- ◆ Selective Acceleration of Critical Sections (SEL)
  - Saturating counters to track false serialization



Fall 2016

Lec.20 - Slide 34

## Outline

- ◆ Background
- ◆ Mechanism
- ◆ **Performance Trade-Offs**
- ◆ Evaluation
- ◆ Related Work and Summary

Fall 2016

Lec.20 - Slide 35

## Performance Tradeoffs


- ◆ **Fewer threads vs. accelerated critical sections**
  - Accelerating critical sections offsets loss in throughput
  - As the number of cores (threads) on chip increase:
    - ▲ Fractional loss in parallel performance decreases
    - ▲ Increased contention for critical sections makes acceleration more beneficial
- ◆ **Overhead of CSCALL/CSDONE vs. better lock locality**
  - ACS avoids “ping-ponging” of locks among caches by keeping them at the large core
- ◆ **More cache misses for private data vs. fewer misses for shared data**

Fall 2016

Lec.20 - Slide 36

## Cache misses for private data

`PriorityHeap.insert(NewSubProblems)`



Fall 2016

Puzzle Benchmark

## Performance Tradeoffs

- ◆ **Fewer threads vs. accelerated critical sections**
  - Accelerating critical sections offsets loss in throughput
  - As the number of cores (threads) on chip increase:
    - ▲ Fractional loss in parallel performance decreases
    - ▲ Increased contention for critical sections makes acceleration more beneficial
- ◆ **Overhead of CSCALL/CSDONE vs. better lock locality**
  - ACS avoids “ping-ponging” of locks among caches by keeping them at the large core
- ◆ **More cache misses for private data vs. fewer misses for shared data**
  - Cache misses reduce if shared data > private data

Fall 2016

Lec.20 - Slide 38


## Outline

- ◆ Background
- ◆ Mechanism
- ◆ Performance Trade-Offs
- ◆ **Evaluation**
- ◆ Related Work and Summary

Fall 2016


Lec.20 - Slide 39

## Experimental Methodology




SCMP

- All small cores
- Conventional locking



ACMP

- One large core (area-equal 4 small cores)
- Conventional locking



ACS

- ACMP with a CSRB
- Accelerates Critical Sections

Fall 2016

Lec.20 - Slide 40

## Experimental Methodology

- ◆ Workloads
  - 12 critical section intensive applications from various domains
  - 7 use coarse-grain locks and 5 use fine-grain locks
- ◆ Simulation parameters:
  - x86 cycle accurate processor simulator
  - Large core: Similar to Pentium-M with 2-way SMT. 2GHz, out-of-order, 128-entry ROB, 4-wide issue, 12-stage
  - Small core: Similar to Pentium 1, 2GHz, in-order, 2-wide issue, 5-stage
  - Private 32 KB L1, private 256KB L2, 8MB shared L3
  - On-chip interconnect: Bi-directional ring

Fall 2016


Lec 20 - Slide 41

## Workloads with Coarse-Grain Locks

Equal-area comparison  
Number of threads = *Best threads*

### Chip Area = 16 cores


SCMP = 16 small cores  
ACMP/ACS = 1 large and 12 small cores



Fall 2016

### Chip Area = 32 small cores

SCMP = 32 small cores  
ACMP/ACS = 1 large and 28 small cores




Lec 20 - Slide 42

## Workloads with Fine-Grain Locks

Equal-area comparison  
Number of threads = *Best threads*

### Chip Area = 16 cores

SCMP = 16 small cores  
ACMP/ACS = 1 large and 12 small cores




Fall 2016

Lec 20 - Slide 43

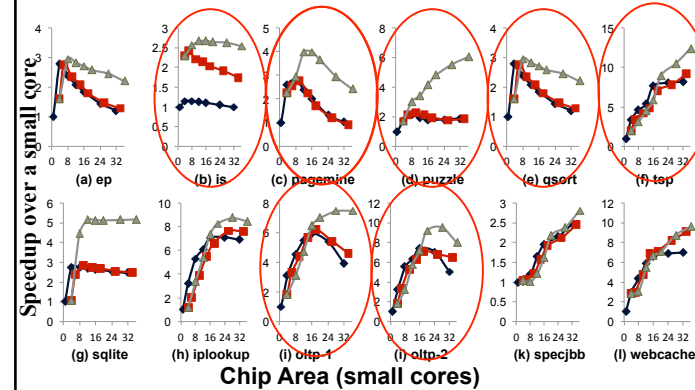
### Chip Area = 32 small cores

SCMP = 32 small cores  
ACMP/ACS = 1 large and 28 small cores



## Equal-Area Comparisons

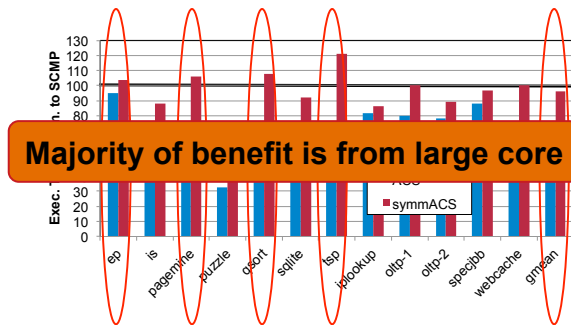
--- SCMP  
--- ACMP  
--- ACS



Fall 2016

Lec 20 - Slide 44

## ACS on Symmetric CMP



Fall 2016

Lec.20 - Slide 45

## Outline

- ◆ Background
- ◆ Mechanism
- ◆ Performance Trade-Offs
- ◆ Evaluation
- ◆ **Related Work and Summary**

Fall 2016

Lec.20 - Slide 46

## Related Work

- ◆ Improving locality of shared data by thread migration and software prefetching (Sridharan+, Trancoso+, Ranganathan+)
  - ACS not only improves locality but also uses a large core to accelerate critical section execution*
- ◆ Asymmetric CMPs (Morad+, Kumar+, Suleman+, Hill+)
  - ACS not only accelerates the Amdahl's bottleneck but also critical sections*
- ◆ Remote procedure calls (Birrell+)
  - ACS is for critical sections among shared memory cores*

Fall 2016

Lec.20 - Slide 47

## Hiding Latency of Critical Sections

- ◆ Transactional memory (Herlihy+)
    - ACS does not require code modification*
  - ◆ Transactional Lock Removal (Rajwar+) and Speculative Synchronization (Martinez+)
    - Hide critical section latency by increasing concurrency
      - ACS reduces latency of each critical section*
    - Overlaps execution of critical sections with no data conflicts
      - ACS accelerates ALL critical sections*
    - Does not improve locality of shared data
      - ACS improves locality of shared data*
- ACS outperforms TLR (Rajwar+) by 18% (details in paper)

Fall 2016

Lec.20 - Slide 48

## Conclusion

---

- ◆ Critical sections limit scalability
- ◆ Accelerate critical sections with a powerful core
- ◆ ACS reduces average execution time by:
  - 34% compared to an equal-area SCMP
  - 23% compared to an equal-area ACMP
- ◆ ACS improves scalability of 7 of the 12 workloads