


ACA

Servers

Fall 2016



Pejman Lotfi-Kamran

Adapted from slides originally developed by Profs. Hill, Hoe, Falsafi and Wenisch of CMU, EPFL, Michigan, Wisconsin

Lec.24 - Slide 1

## Where Are We?

Fr	Sa	Su	Mo	Tu
	27-Shahrivar		29-Shahrivar	
	3-Mehr		5-Mehr	
	10-Mehr		12-Mehr	
	17-Mehr		19-Mehr	
	24-Mehr		26-Mehr	
	1-Aban		3-Aban	
	8-Aban		10-Aban	
	15-Aban		17-Aban	
	22-Aban		24-Aban	
	29-Aban		1-Azar	
	6-Azar		8-Azar	
	13-Azar		15-Azar	
	20-Azar		22-Azar	
	27-Azar		29-Azar	
	4-Dey		6-Dey	

Fall 2016

Lec.24 - Slide 2

- This Lecture  
– Server

- Next Lecture:  
– Data Center

## What do Scale-Out Apps Need? [ISCA 2012] [IEEE Micro 2012]

Cores share instructions


- Large code footprint fits in LLC

Data is in memory

- Data footprint dwarfs LLC

Cores don't communicate

- Independent requests
- Mostly Read-only accesses



Fall 2016

Poor match to existing server processors

Lec.24 - Slide 3


## Conventional Processors: Major Mismatch

✗ Few fat cores

- Limited on-chip parallelism

✗ Large LLC

- Dwarfed by vast datasets
- Slow to access
- Takes up much die area



e.g., Intel Westmere

Large cores & LLC wasteful on scale-out apps

## Small-chip Processors: Not Enough

### ✗ Few lean cores

- Low parallelism



e.g., Calxeda

### ✗ Modestly-sized cache

- Fast to access, but...
- takes up large fraction of die area

Fall 2016

Lec.24 - Slide 5


## Emerging Tiled Processors: Not Optimal

### ✓ Many lean cores

- High parallelism

### ✗ Large distributed cache

- Dwarfed by vast datasets
- Slow to access
- Takes up much die area



e.g., Tilera Tile64

### ✗ Tiled organization

- Distance hurts performance

**Large LLC & distance prevent scalability**

Fall 2016

Lec.24 - Slide 6

## Outline

- Why not today's processors?
- What do scale-out workloads need?
- Scale-Out Processors
  - Overview
  - Design considerations

Fall 2016

Lec.24 - Slide 7

## Roadmap for Scale-Out Processor

1. Eliminate wasteful capacity in the LLC
  - Shrink the cache, add cores in freed space
2. Reduce delay to LLC
  - Partition the die into independent multi-core *Pods*
3. Enable technology scalability
  - Do not interconnect the pods


Fall 2016

Lec.24 - Slide 8

## Step 1: Less cache, more cores

### ✓ Small LLC

- Capture instructions
- More area for the cores



Fall 2016


Lec.24 - Slide 9

## Step 1: Less cache, more cores

### ✓ Small LLC

- Capture instructions
- More area for the cores

### ✓ More cores for higher throughput



Fall 2016


Lec.24 - Slide 10

## Step 2: Form “pods” to reduce distance

### ✓ Small LLC

### ✓ More cores for higher throughput

### ✗ Fast path to LLC for instructions



Fall 2016

Lec.24 - Slide 11

How to choose the optimal pod?

## Sizing Pods: Not So Simple

Too few cores → limited parallelism

Too many cores → long distance to LLC

- Slower instruction fetch


Cache dominates  
die area



Few cores



Distance hurts performance




Many cores

Fall 2016

Lec.24 - Slide 12

How do we characterize optimality?

## Our Optimization Criterion: Performance Density (PD)



Fall 2016

PD pinpoints optimal use of silicon

Lec.24 - Slide 13

## Step 3: Scale-Out Processors


One or more pods

Each pod is a standalone server

- Runs a full software stack

No inter-pod connectivity or coherence

- Enhances scalability



Fall 2016

Inherently optimal & scalable

Lec.24 - Slide 14

## Methodology

Flexus simulation infrastructure [Wenisch '06]

CMP analytic model [Hardavellas '09]

TCO model [Hardy '11]

### Core Parameters

- Out-of-order, 3-way, 2 GHz
- L1 (I & D): 32KB, 2-way

### Chip budget

- Die area: 280 mm²
- Power: 95W
- BW: 6 x 3.2 GT/s DDR4

### Chip components

- Technology node: 20 nm
- Core: 1.1 mm²
- 1MB cache: 1.2 mm²
- Mem channel: 12 mm²


### Datacenter

- Power: 20 MW
- Rack: 42 U, 17 kW
- Server: 64 GB, N sockets

Fall 2016

Lec.24 - Slide 15

## Pod Design Space Exploration




Fall 2016

Optimal Pod: 32 cores & 4MB of cache

Lec.24 - Slide 16

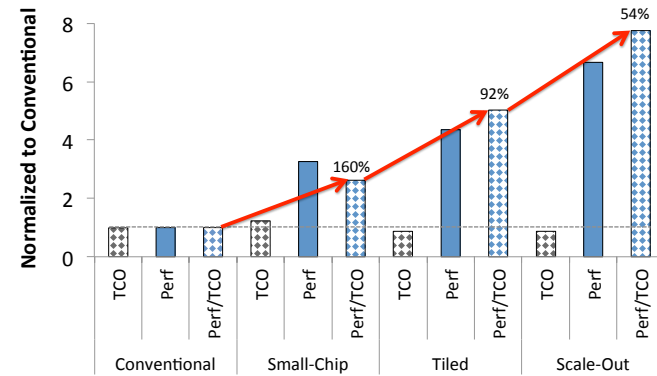
## Processor Efficiency



SOP: highest performance per mm<sup>2</sup>

Lec.24 - Slide 17

## Datacenter Efficiency



SOP: highest performance/TCO

Lec.24 - Slide 18

## Outline

- Why not today's processors?
- What do scale-out workloads need?
- Scale-Out Processors
  - Overview
  - Design considerations

Fall 2016

Lec.24 - Slide 19


## How do we architect a pod? [MICRO 2012]

### Pod characteristics:

- Several MB of LLC
- Many cores

### Challenge:

- Efficiently connect the pieces
- Maximize PD



Fall 2016

Lec.24 - Slide 20


## Tiled Pod

### Mesh NOC

- Nearest-neighbor connectivity
- 5-ported routers

### Distributed LLC

- ✓ Low complexity
- ✓ Small NOC area
- ✗ High LLC access latency




Low cost, low performance

Fall 2016

Lec.24 - Slide 21

## Performance: Tiled Pod vs Ideal



Fall 2016

Lec.24 - Slide 22


## Fast Tiled Pod

### Flattened Butterfly NOC Kim, MICRO '07

- Rich connectivity
- Many-ported routers

### Distributed LLC

- ✗ High complexity
- ✗ Large NOC area (7x of mesh!)
- ✓ Low LLC access latency




High cost, high performance

Fall 2016

Lec.24 - Slide 23

## Off-the-shelf designs not ideal



Exploit workload characteristics for efficiency


## How do threads access data?

Where is the code and data?

- Code: in LLC
- Data: in memory
- Nothing useful in remote L1s!

Data access pattern: *bilateral*

- L1 ↔ LLC
- L1 ✗ L1




Fall 2016

Lec.24 - Slide 25

## What is the actual traffic pattern?

Distributed LLC: all-to-all traffic

- Mismatch with underlying bilateral access pattern
- Requires expensive connectivity for low latency



Turn all-to-all traffic into bilateral to reduce cost!

## Roadmap for an “Ideal” Pod

Starting point: fast tiled pod


1. LLC in the center
  - Decouple cores and LLC banks
2. Remove links from flattened butterfly
  - Leverage the *bilateral* traffic pattern to lower cost
3. Share and specialize core-to-LLC interconnect

Fall 2016

Lec.24 - Slide 27

## Step 1: LLC in the Center

- ✓ Decouple core and LLC tiles
  - Natural fit for the bilateral traffic pattern




Fall 2016

Lec.24 - Slide 28

## Step 2: Remove unneeded links

- ✓ Core-to-core connectivity not needed
  - Less cost, same performance
- ✓ Rich intra-LLC connectivity helps performance
  - Expense limited to a fraction of the die




Fall 2016

Lec.24 - Slide 29

## Step 3: Share links and specialize

- ✓ Cores share links to LLC
  - Dedicated core-to-LLC links have poor cost/perf
- ✓ Specialize *request/reply* (to/from LLC) networks
  - Maximize cost/perf




Fall 2016

Lec.24 - Slide 30

## NOC-Out: near-ideal pod design

### Request network

- Tree topology
- Each node: 2-to-1 flow-controlled mux




Fall 2016

Lec.24 - Slide 31

## NOC-Out: near-ideal pod design

### Reply network

- Tree topology
- Each node: 1-to-2 flow-controlled demux



Fall 2016

Lec.24 - Slide 32




## NOC-Out: near-ideal pod design

LLC network: flattened butterfly topology

- Expense limited to a fraction of the die

*Request & Reply* networks: tree topology


- Limited connectivity for efficiency
- SW stack unaffected



Fall 2016

Lec.24 - Slide 33

## Area Breakdown



**NOC-Out: FBfly's performance at 1/10<sup>th</sup> cost**

## Summary

### Scaling Trends

- More lean cores
- Beware of larger caches -> slow you down

### Scale-out workloads

- Lots of request-level parallelism
- Large instruction footprints
- Little core-to-core communication

### Chip organization & design

- Can improve Performance/TCO

Fall 2016

Lec.24 - Slide 35